



K8s Networking

INTRODUCTION GUIDE

September 13, 2022



Objectives and Agenda

K8s networking 101

OBJECTIVES

1. K8S COMPONENTS AND ARCHITECTURE OVERVIEW
2. K8S NETWORKING COMPONENTS OVERVIEW

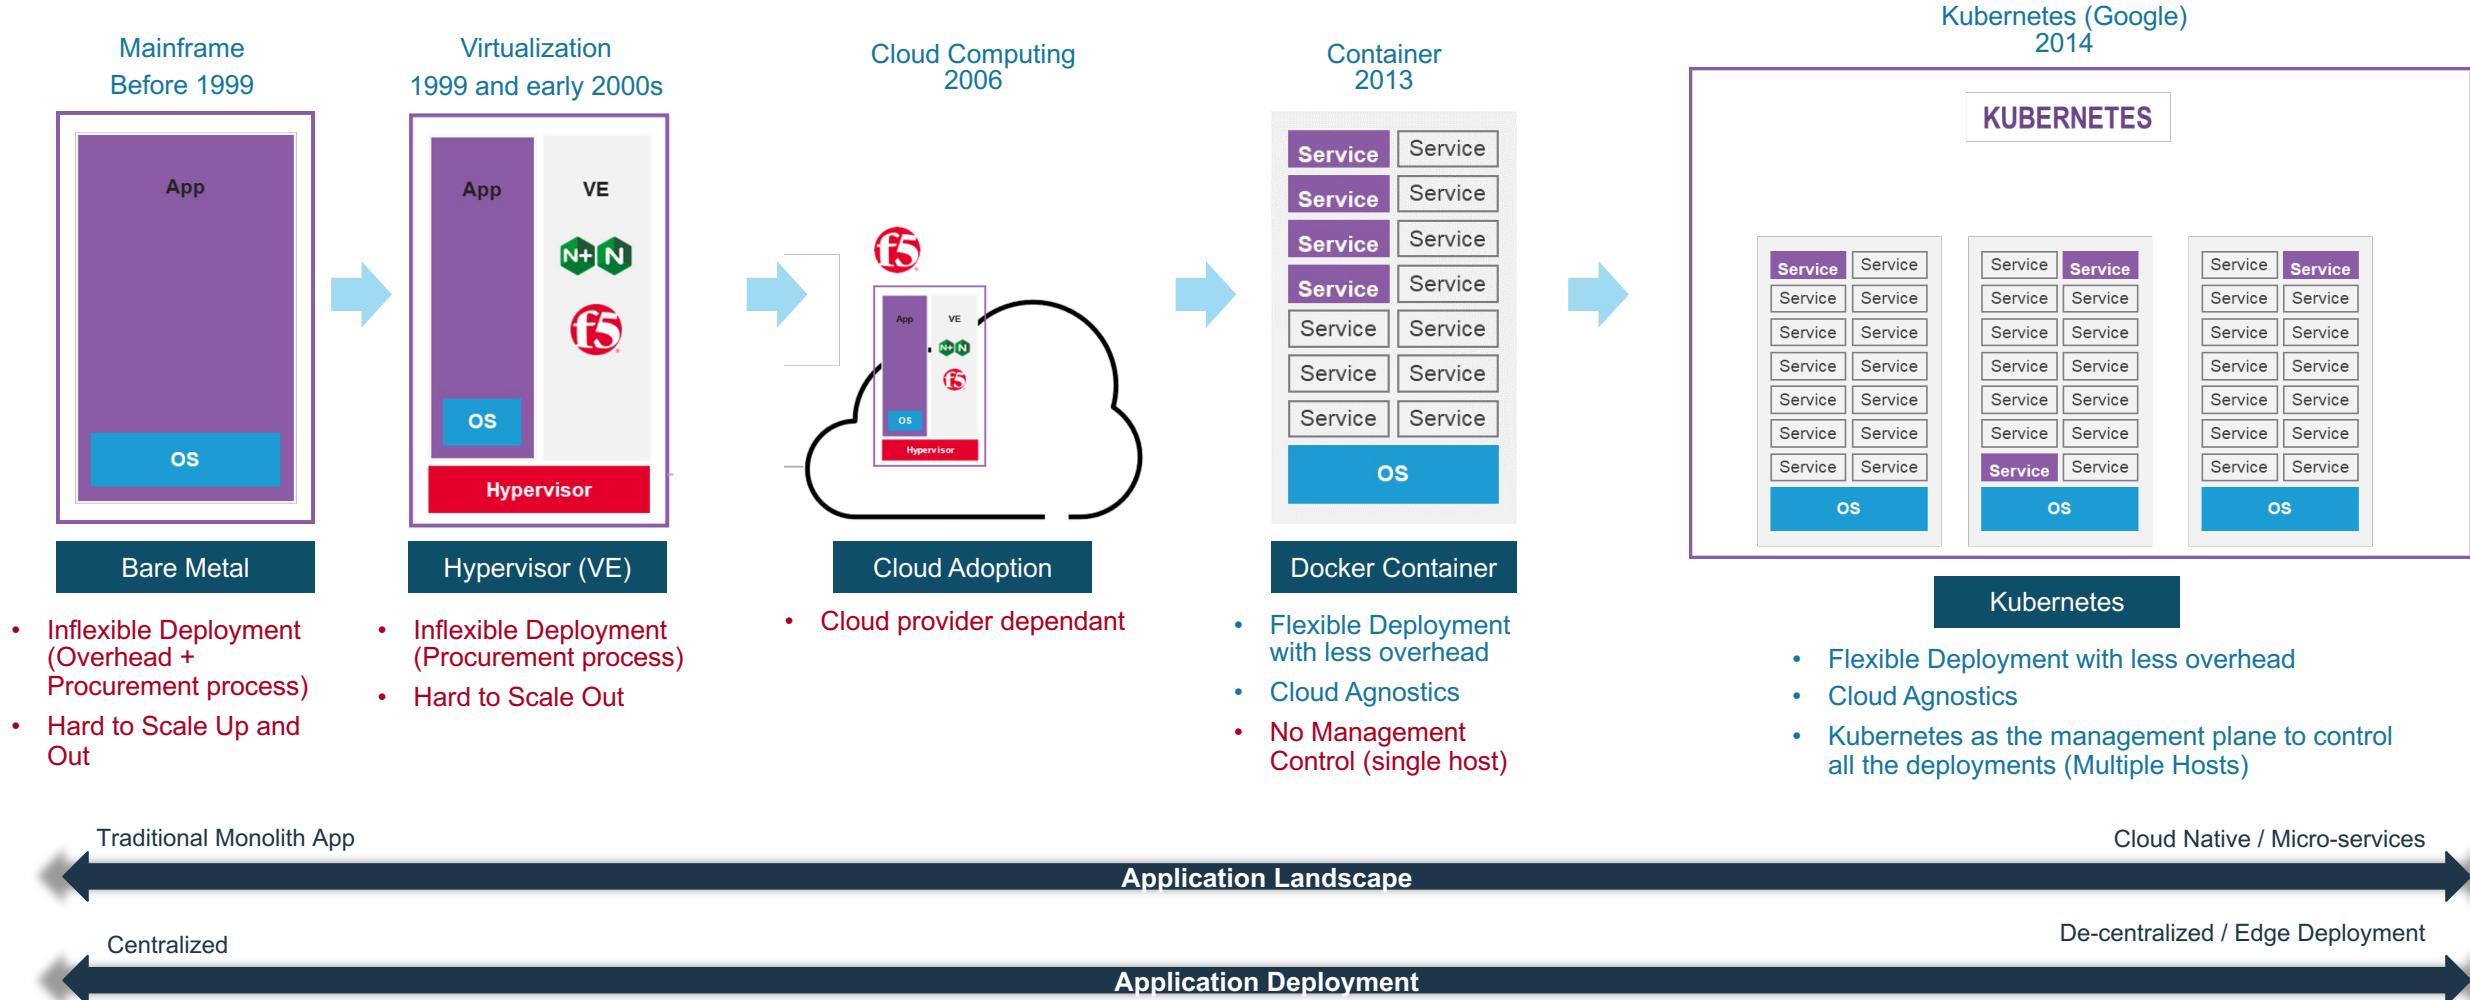
AGENDA

1. MODERN APPS JOURNEY
 2. K8S ARCHITECTURE AND COMPONENTS
 3. K8S CONCEPT
 - kubectl & kubeconfig (Lab 1.1)
 - YAML File and Pods (Lab 1.2)
 - Deployments (Lab 1.3)
 - Namespaces (Lab 1.4)
 - Services (Lab 1.5)
 4. SUMMARY
 5. SECTION QUIZ
- CNI (Lab 1.6)
 - DNS (Lab 1.7)
 - ClusterIP (Lab 1.8)
 - NodePort (Lab 1.9)

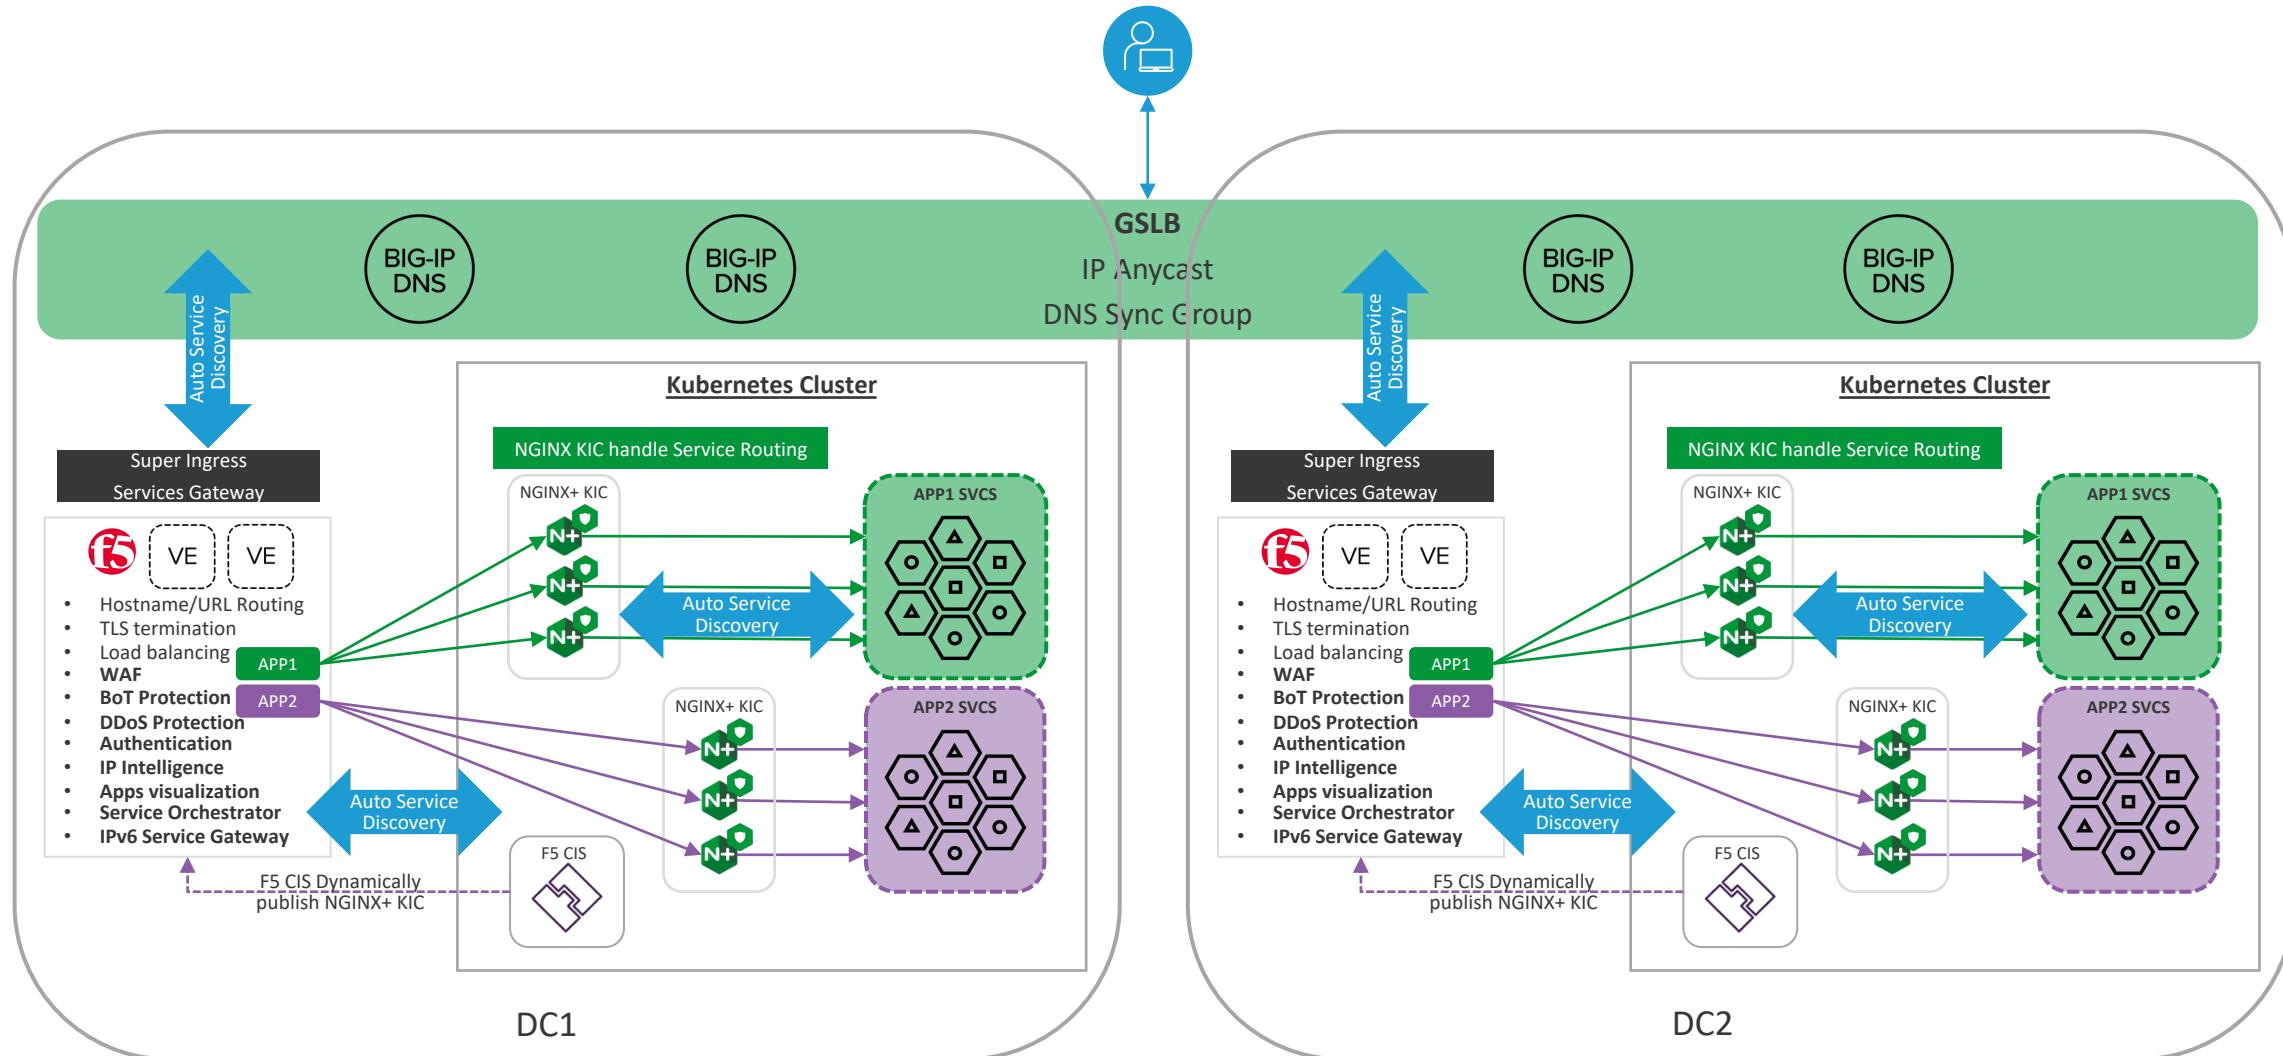
MODERN APPLICATION JOURNEY

Modern Apps/DevSecOps journey

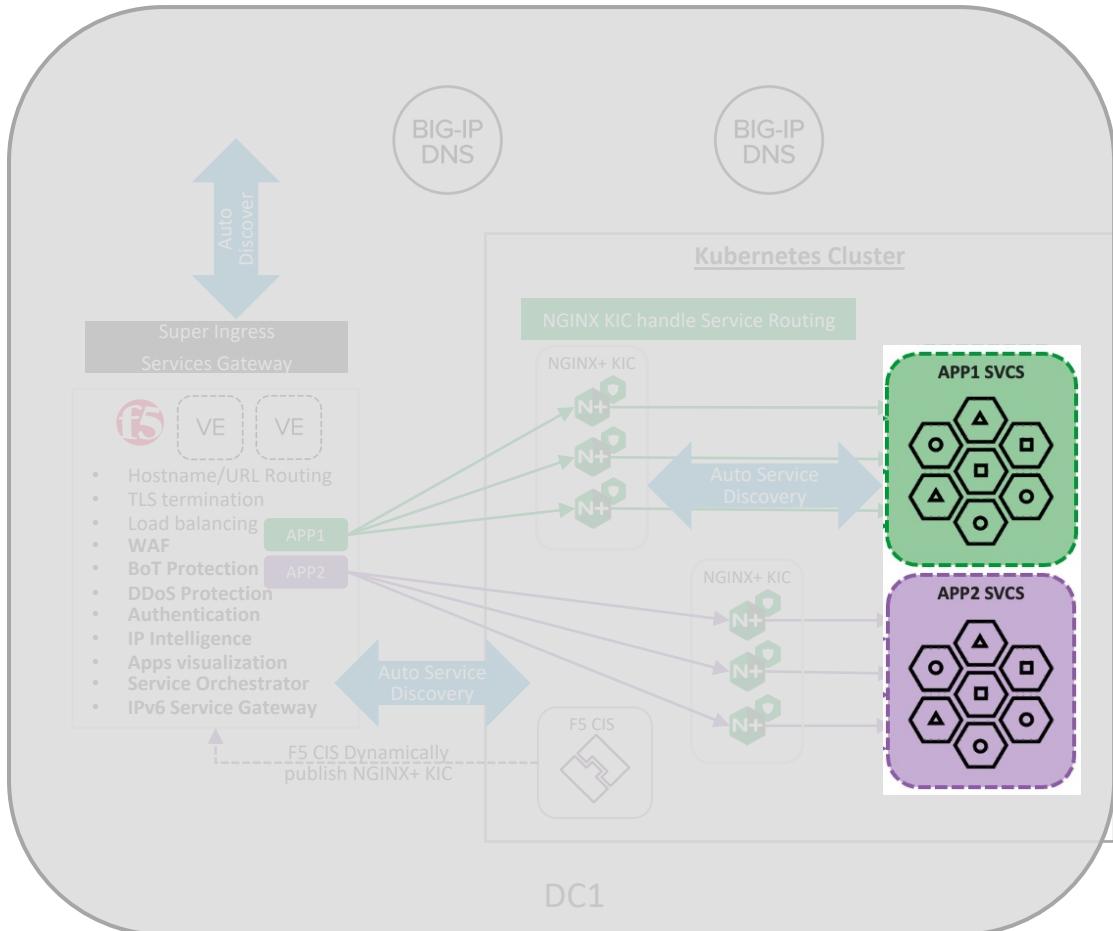
THE EVOLVEMENT OF APPLICATIONS



Agenda : End-goal with Active-Active K8s Clusters Architecture



Day 1 Agenda



Chapter 1 : Modern Applications and DevOps Basics (Theory + Lab)

- Legacy vs VM vs Container (Microservices example in Real Life)
- Getting comfortable with Command Line Editors
 - SSH Commands
 - Using VI Editor (Delete entire line, save, undo, etc)
 - Visual Studio (e.g., F5/NGINX Plug-in)

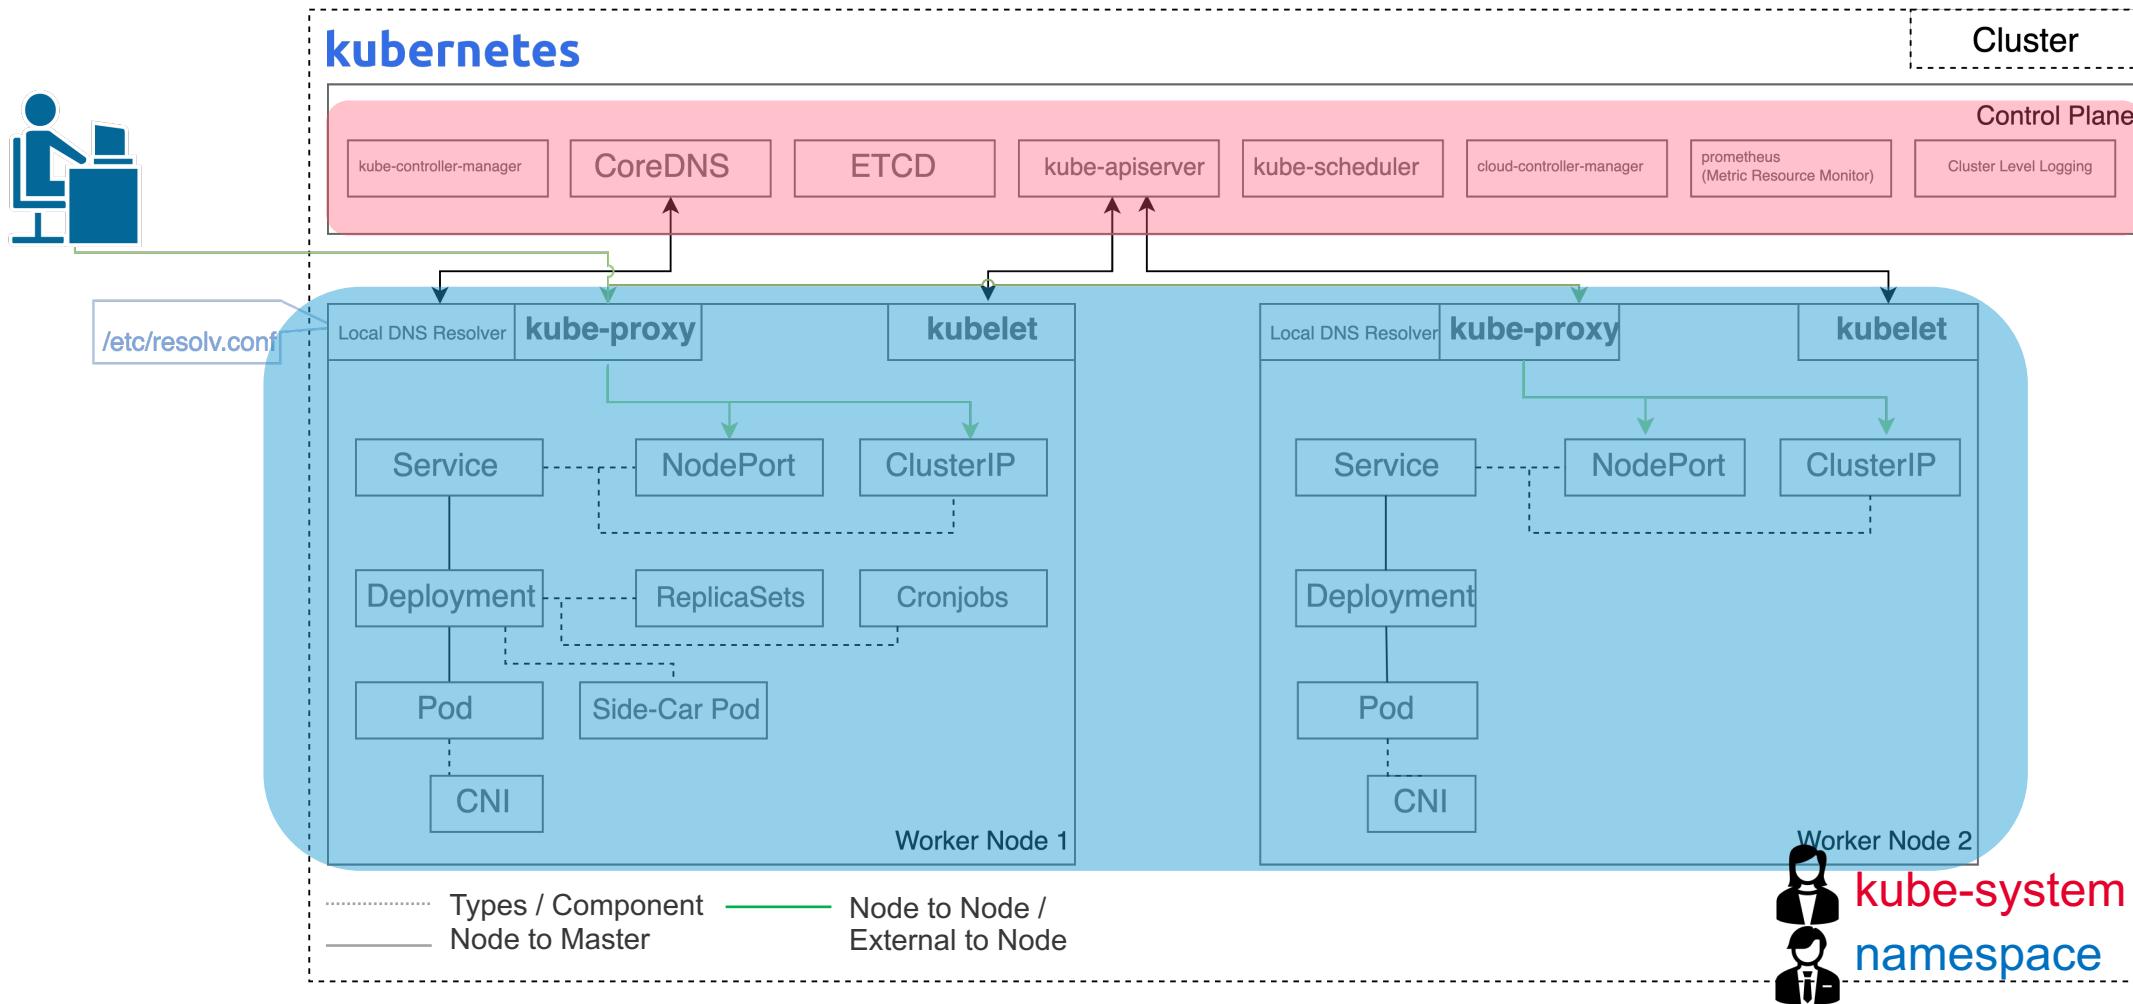
Chapter2 : Kubernetes for Beginners (Theory + K8s Lab)

- Containers and Pods
- YAML Basics – How to read K8s manifests
- K8s Deployment
- K8s Networking: ClusterIP, Services, LB, NodePort

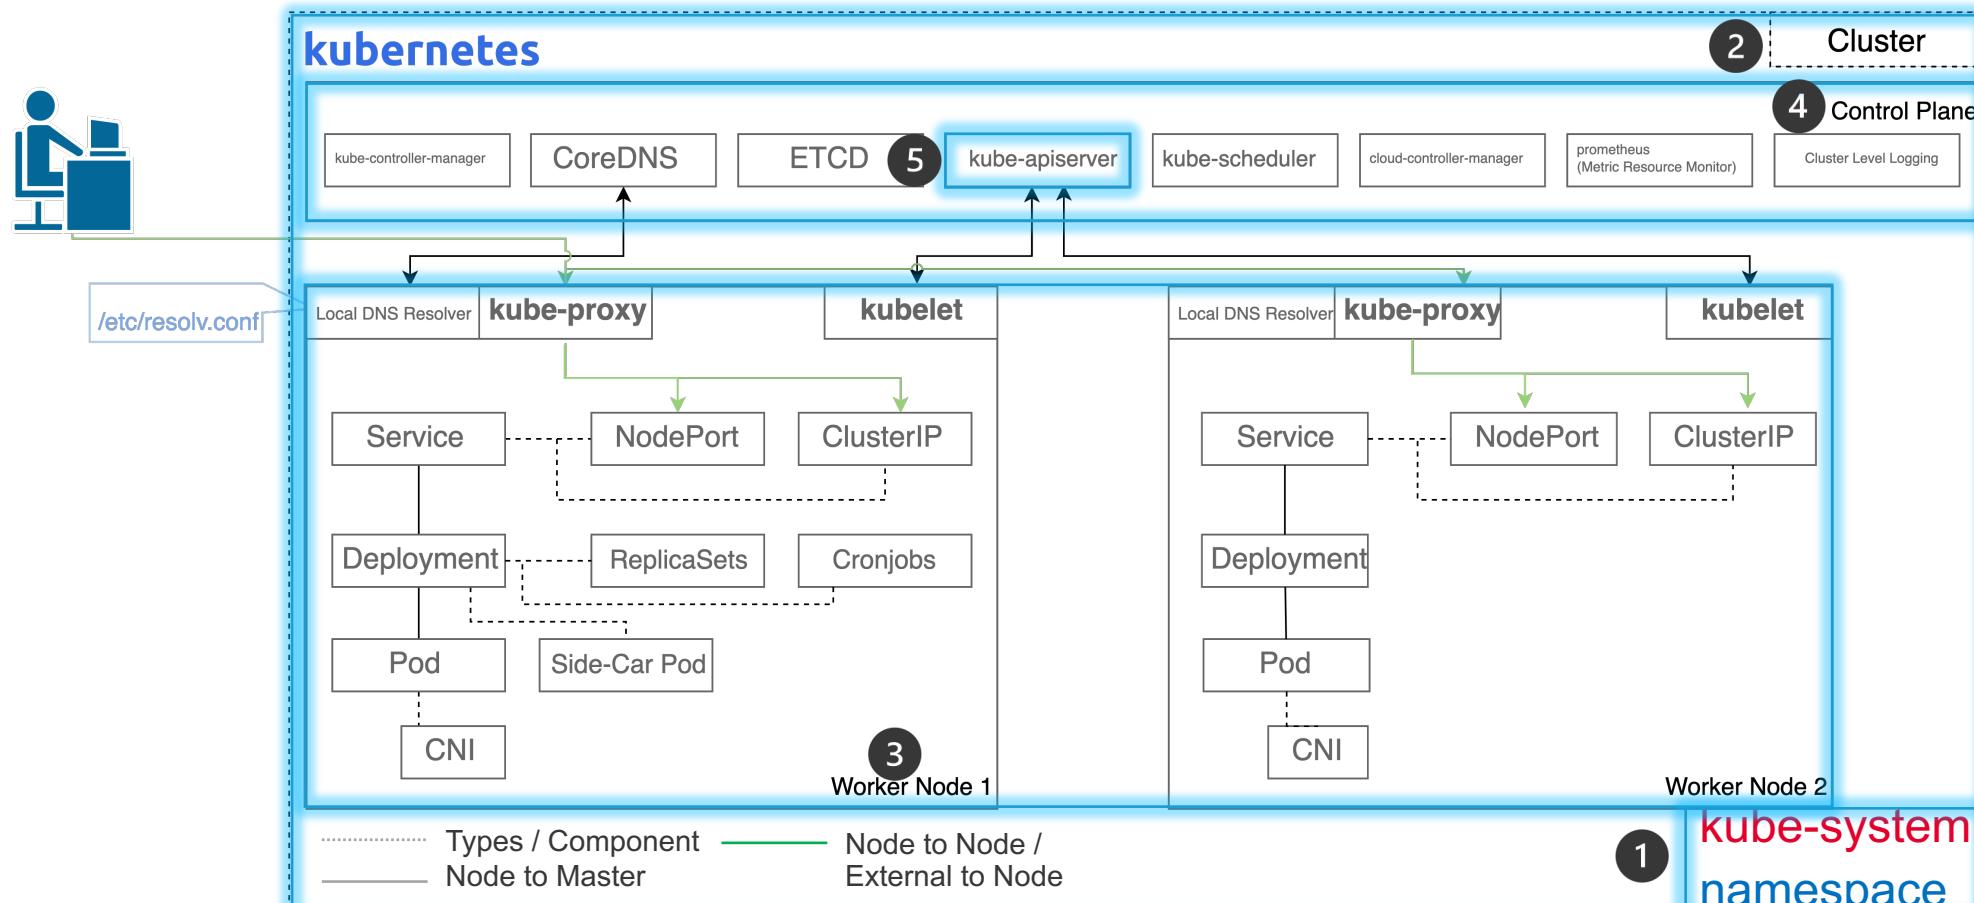
Chapter3: K8s and NGINX Ingress Controller Integration (Theory + Lab)

- What is Ingress Controller?
- How does NGINX+ look like as Ingress Controller?
- NGINX+ value, strength and positioning

Typical K8s deployment for microservices



K8s Topology: Cluster Overview & Master Node



1 Namespace: A way to divide cluster resources between multiple users

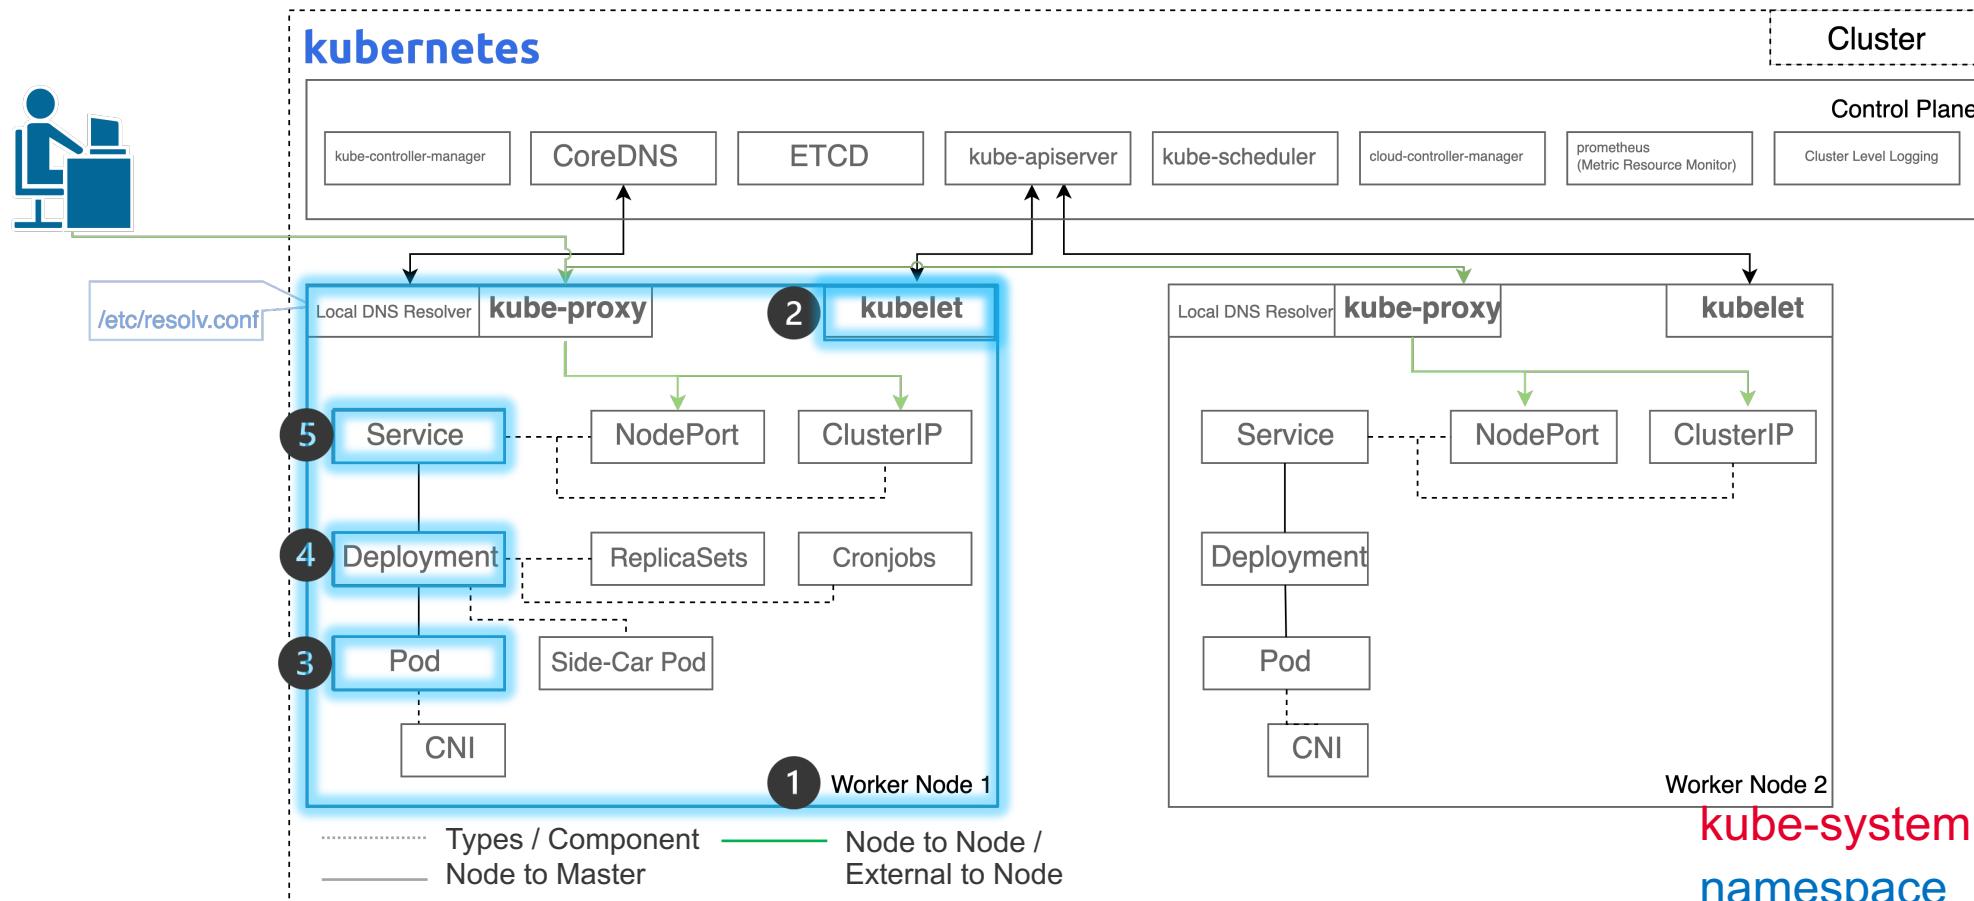
2 Cluster: Pool of Nodes working together intelligently

3 Node: Representation of a virtual machine

4 Control Plane (Master Node): responsible for cluster management like workload distribution and state of the cluster

5 kube-apiserver: API Gateway for operators to the cluster. Used to manipulate each node's components via the usage of API.

K8s Topology: Worker Node



1 Worker Node: runs the workloads for the master node.

2 Kubelet: is an agent that helps the master node to ensure the pods are running within a node

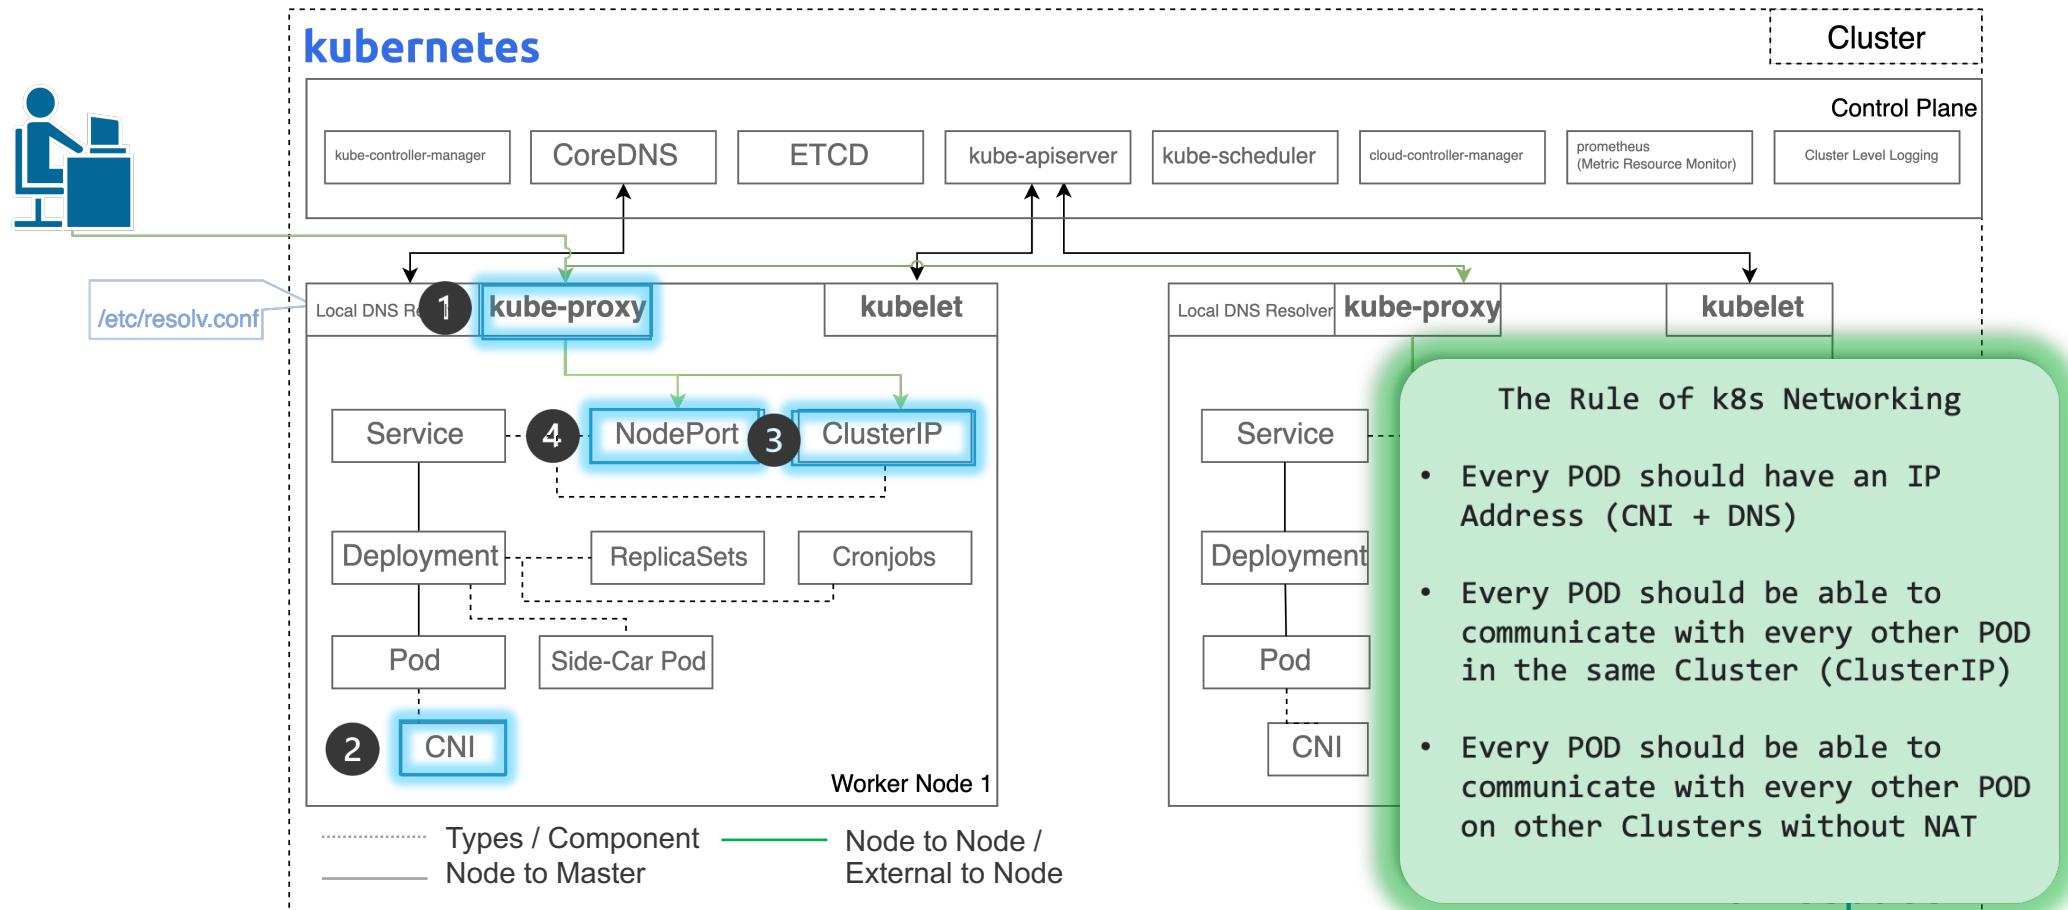
3 Pod: runs the application itself in the form of containers

4 Deployment: used as instruction to inform the Kubelet how to create or modify pods in a controlled manner

5 Service: used to expose a set of deployments to be consumed by other pods within or outside the cluster. (More on that in the next section)



K8s Topology: Network Connectivity



1 **kube-proxy:** a network proxy that runs on each node in your cluster. It maintains network rules on nodes.

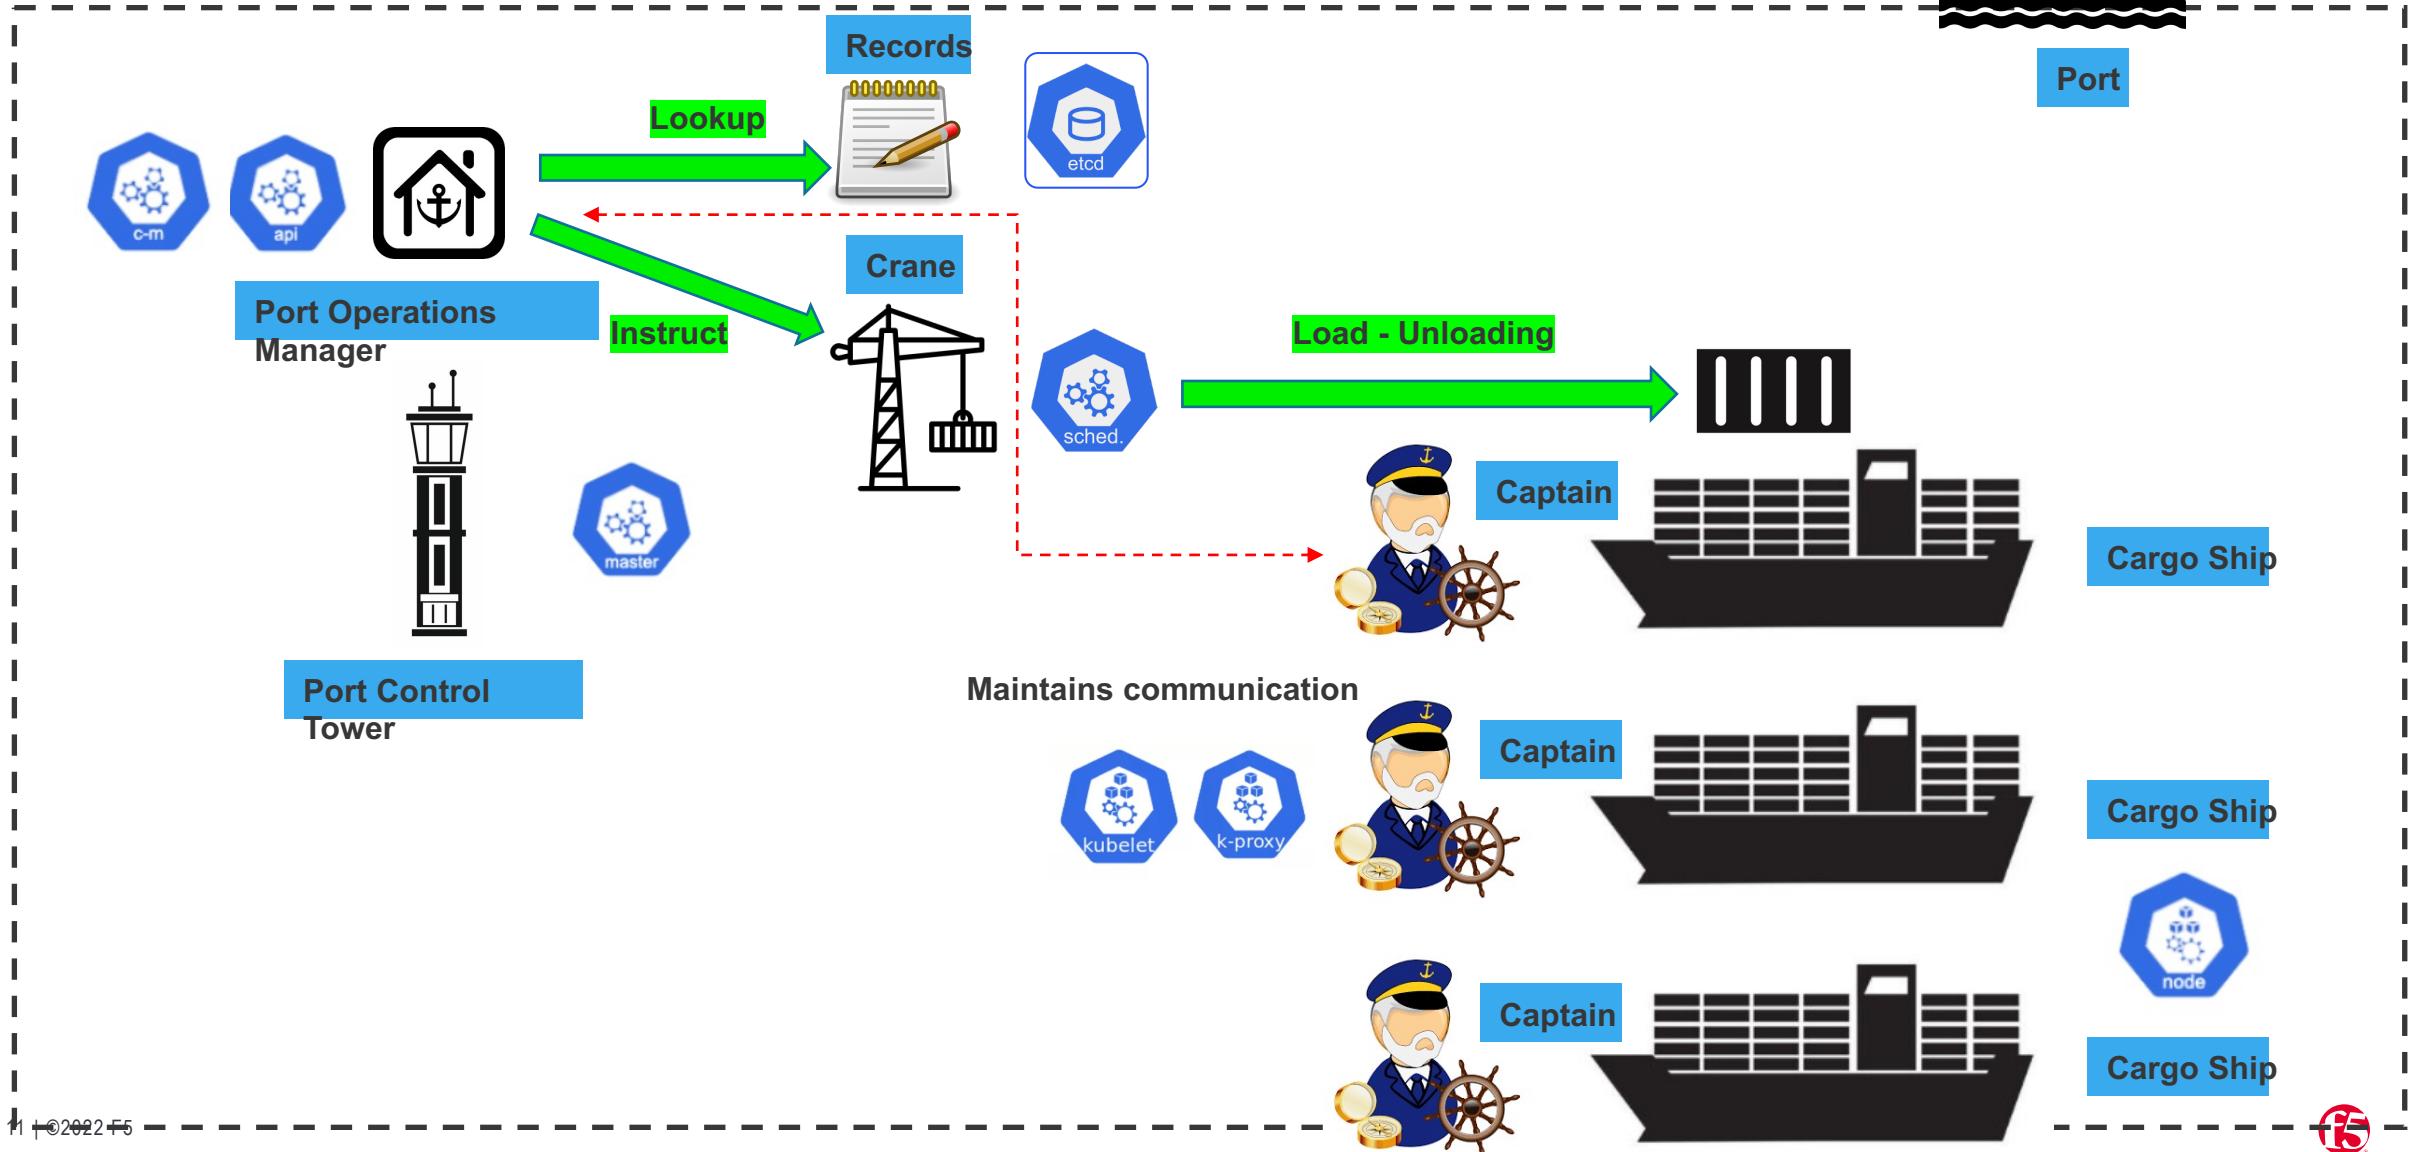
2 **Container Network Interface:** Allows the node to be identified in the network

3 **ClusterIP:** Communication of pods in the same cluster

4 **NodePort:** Exposing Cluster to external

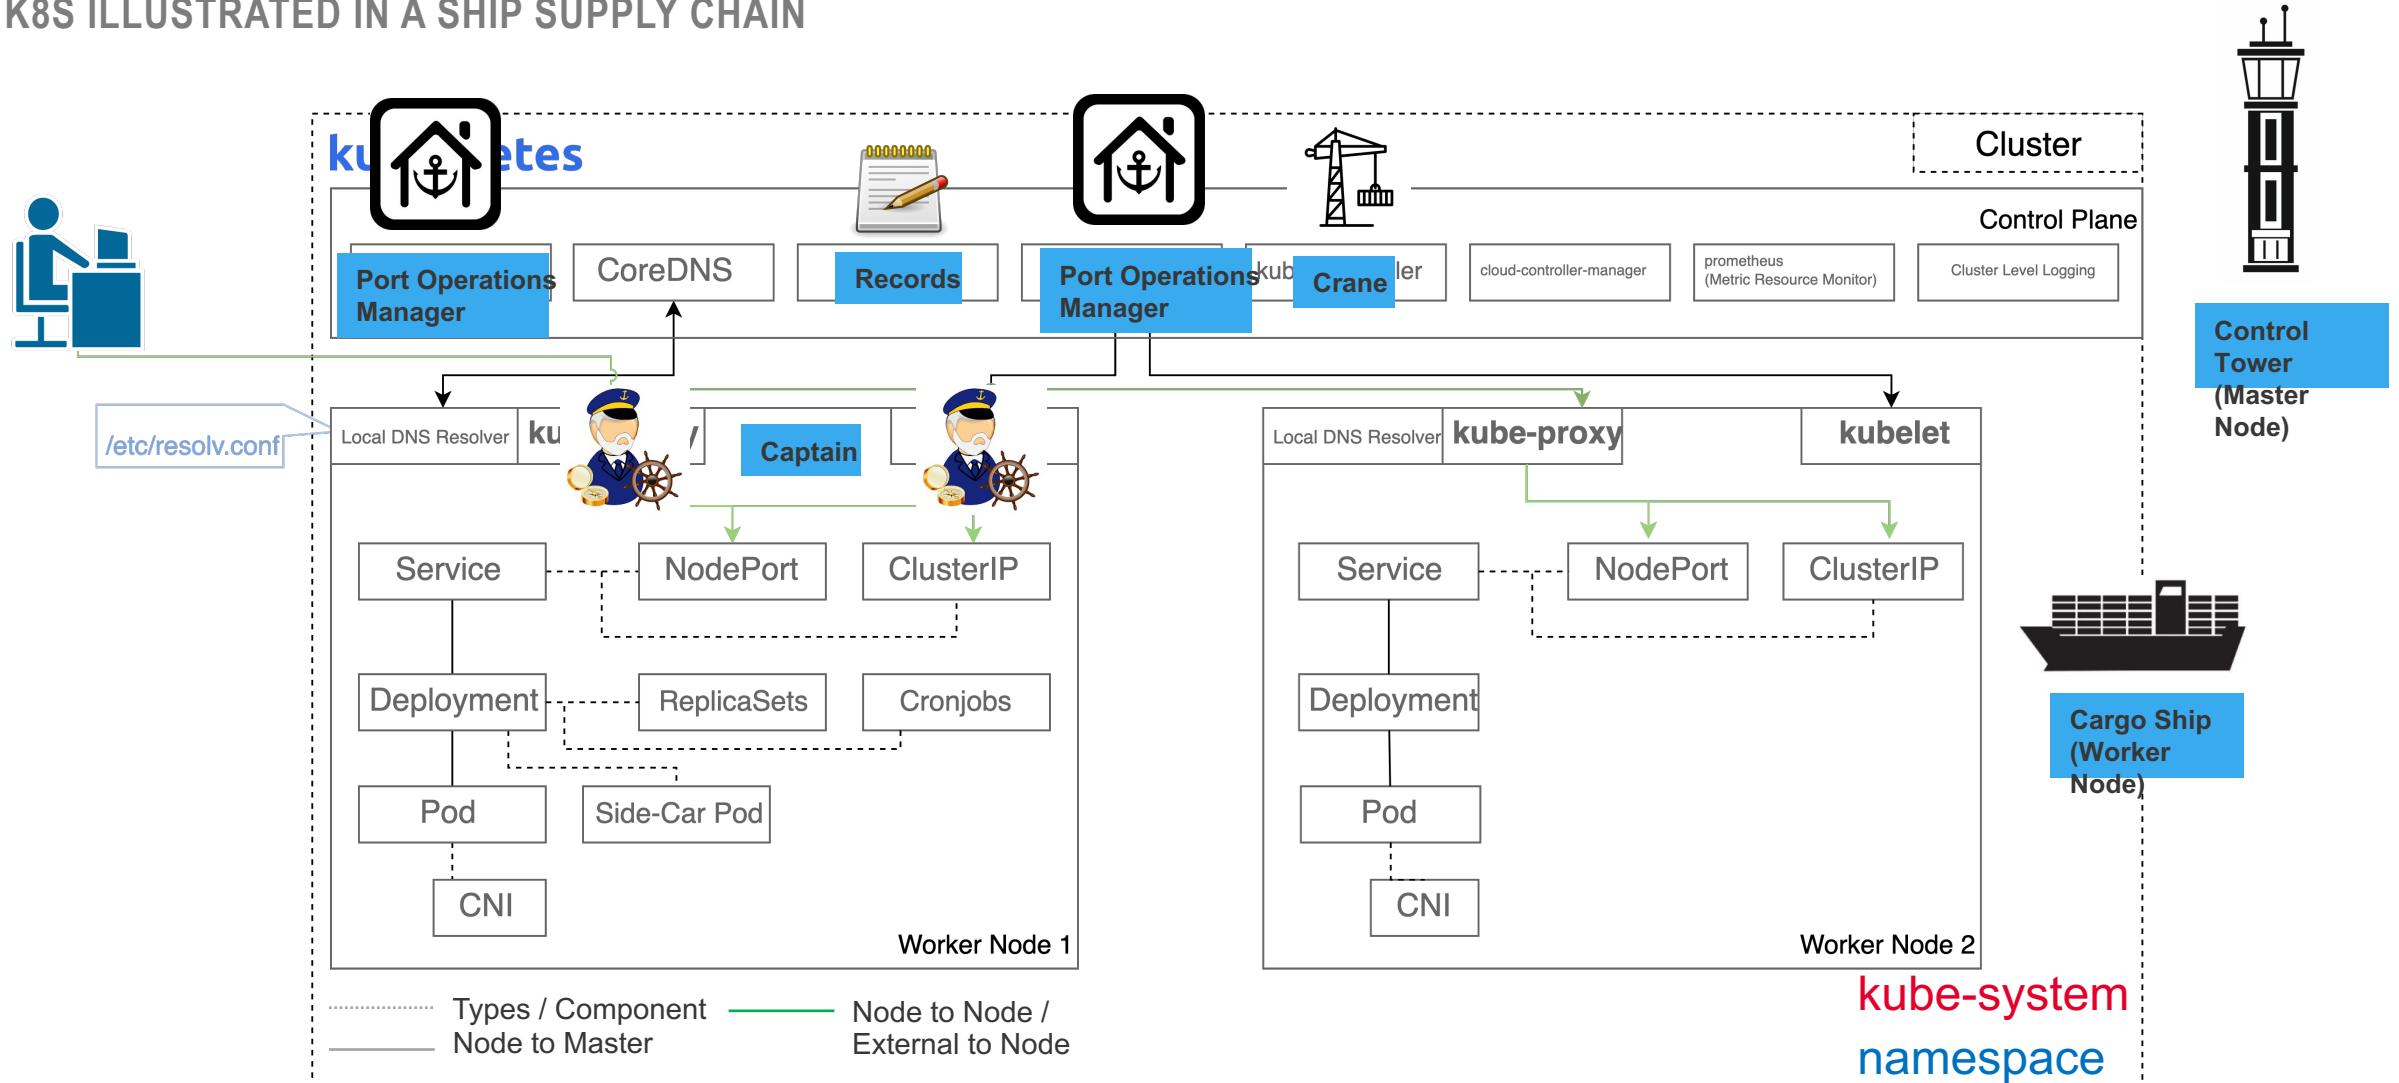
K8s Illustrated

K8S ILLUSTRATED IN A SHIP SUPPLY CHAIN

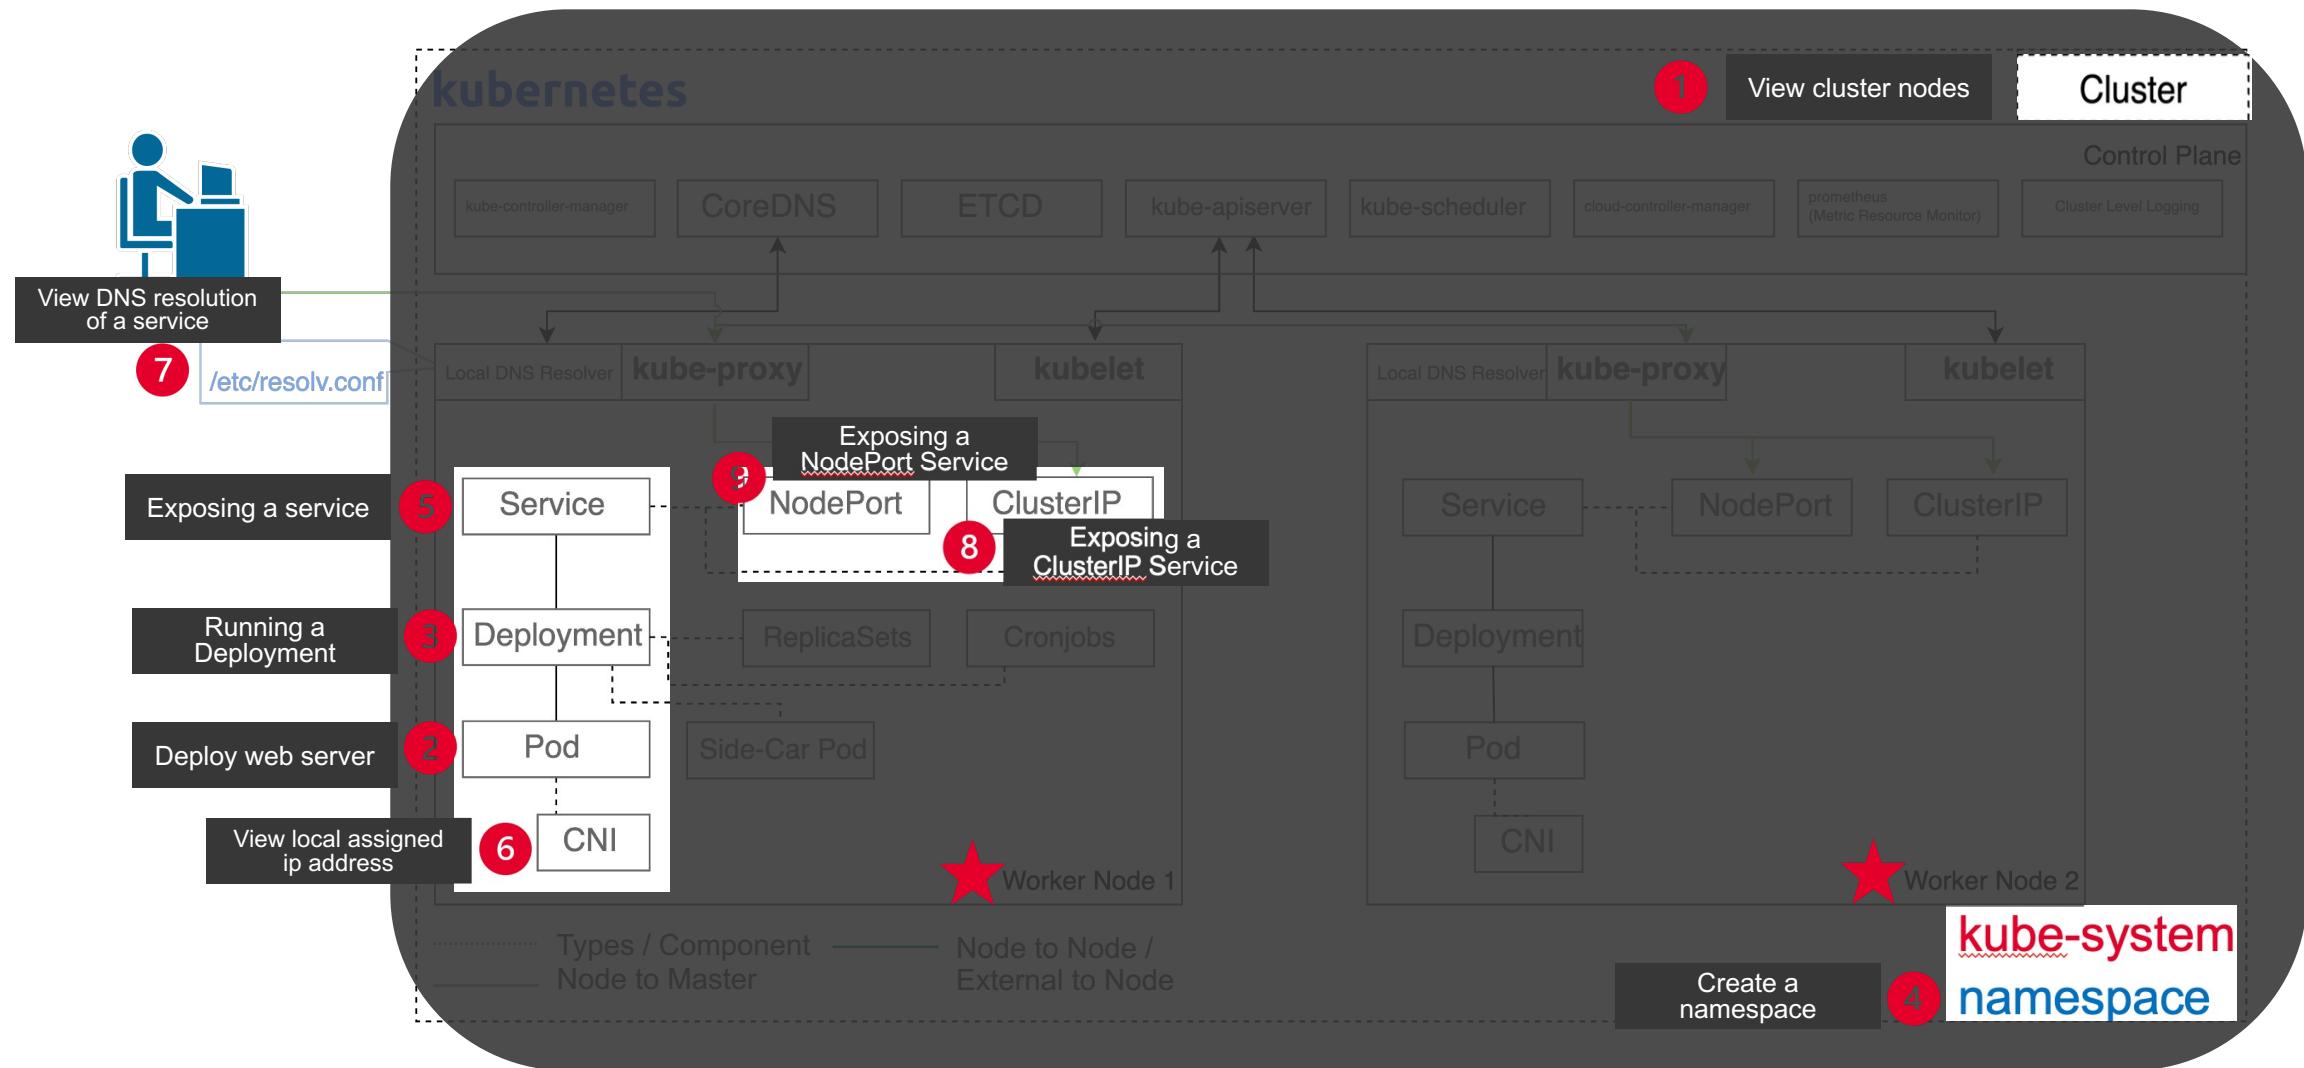


K8s Illustrated

K8S ILLUSTRATED IN A SHIP SUPPLY CHAIN



Lab Task Overview





Lab Guide Introduction and spinning of labs

10 Mins

Kubernetes Concepts

kubectl and kubeconfig

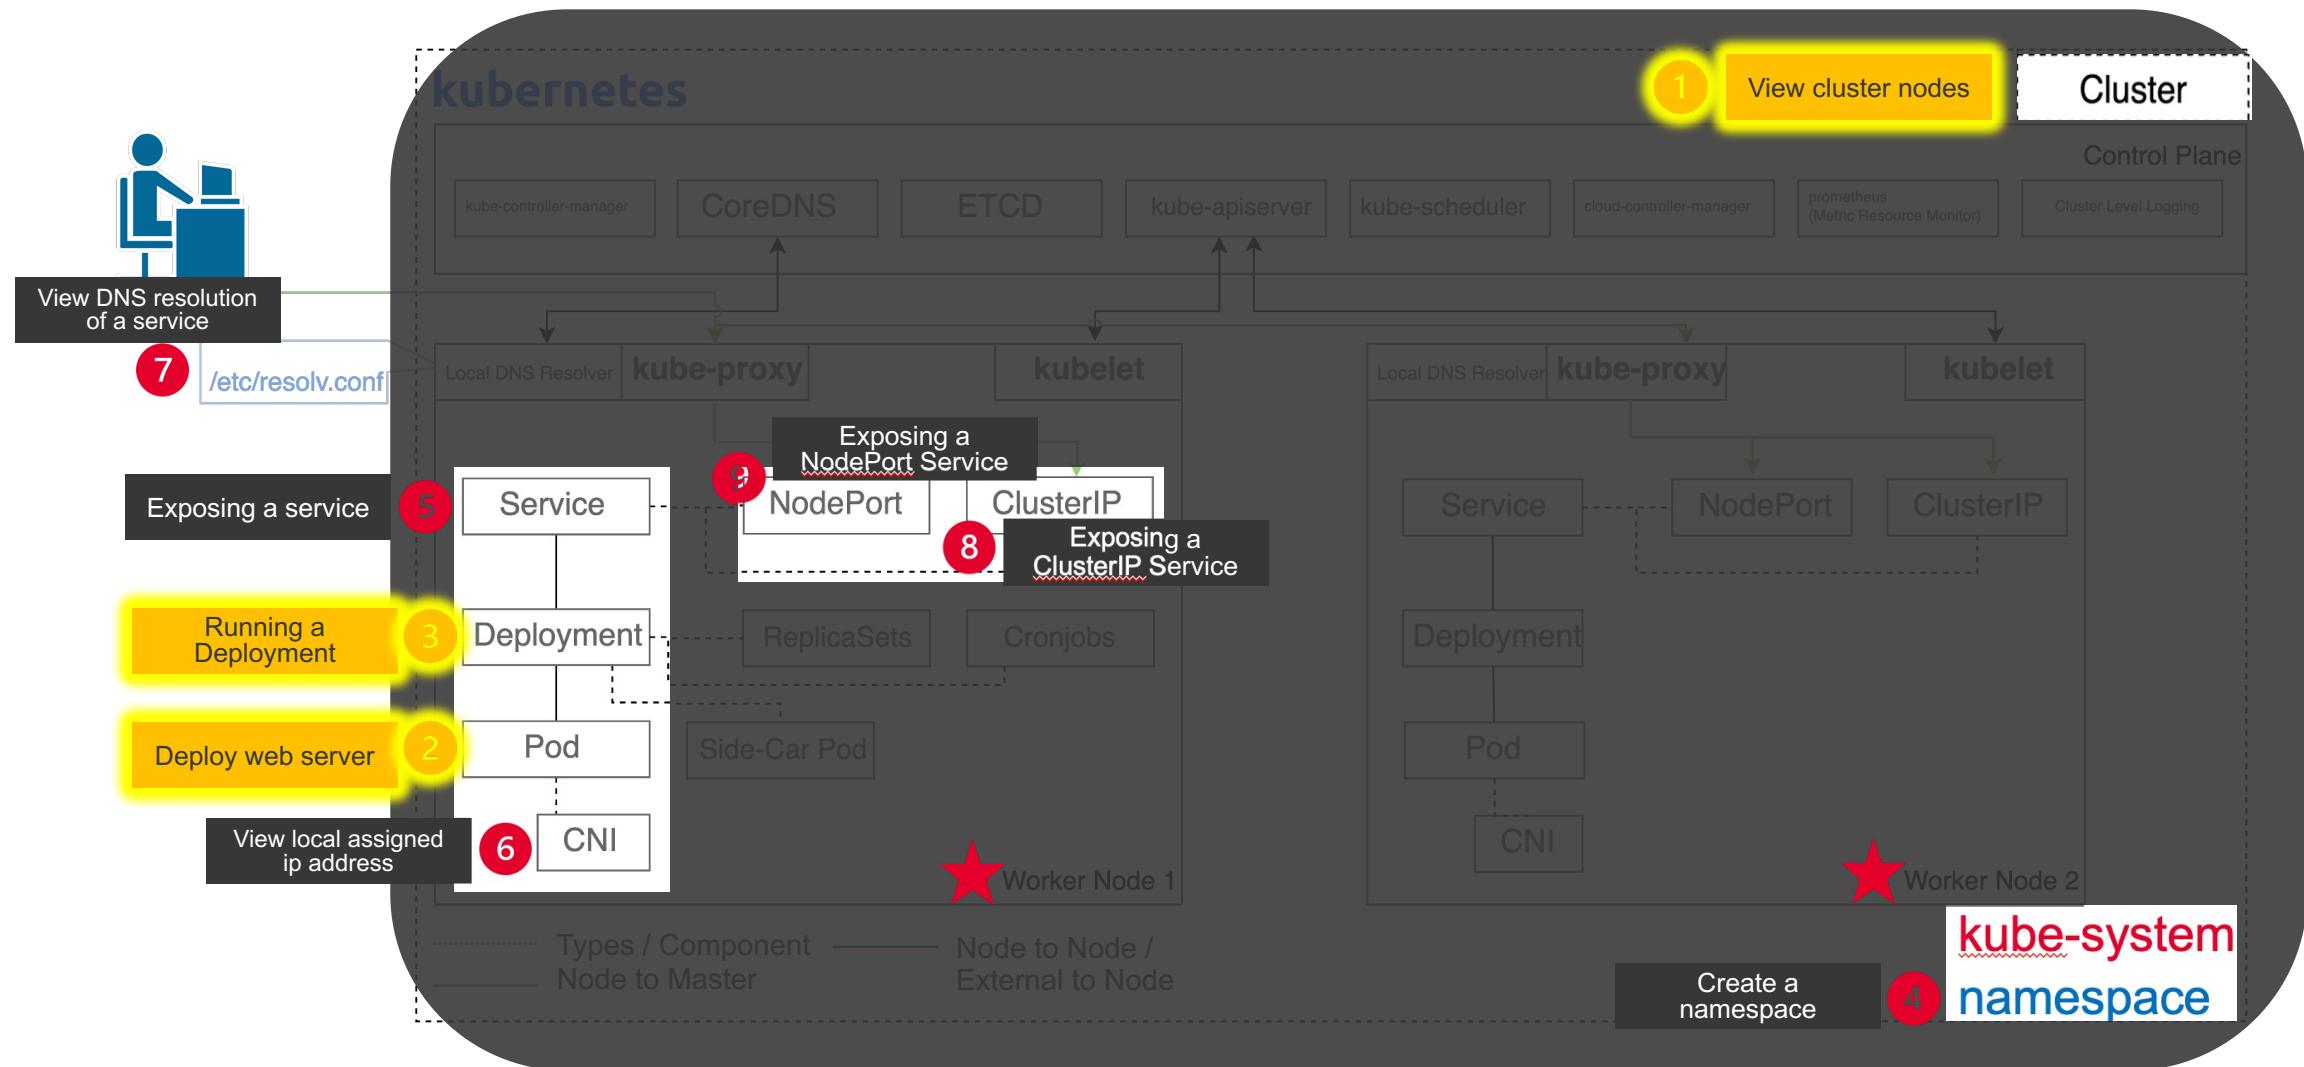
yaml files and pods

deployments

namespaces

services

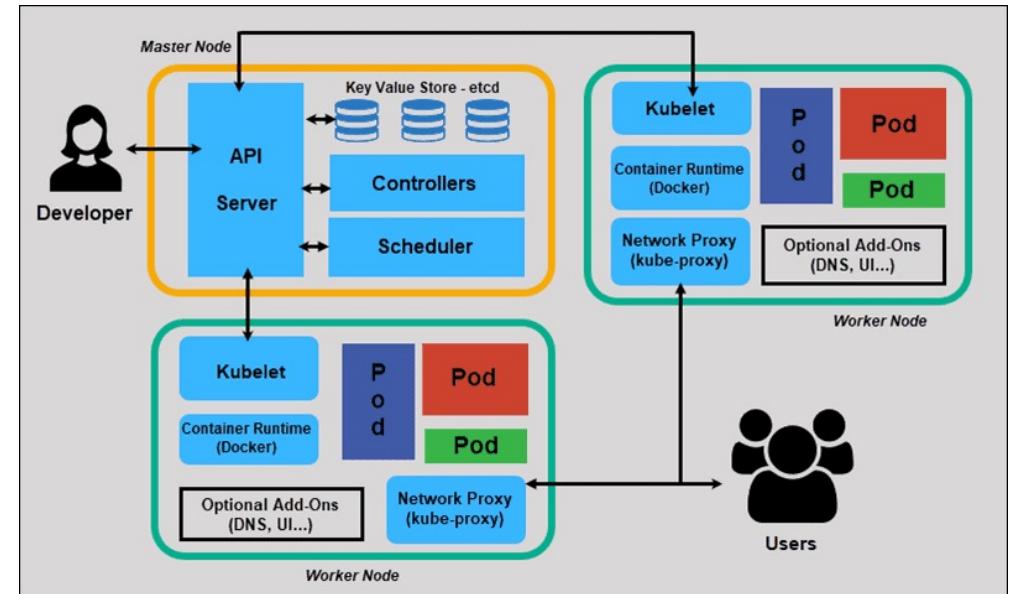
Lab Task Overview



K8s Concept – Lab 1.1

KUBECTL AND KUBECONFIG

- **Kubectl** is the command for developer to make a call to API Server (Master Node)
- In the environment, where K8s master nodes are more than one, the API endpoint can be fronted by a load-balancer to load-balance between the master nodes
- Commands to test the API call and Environment,
 - **kubectl get nodes** ~> List of all nodes
 - Use option **-o wide** to give more information
 - **kubectl get nodes -o wide**



```
server: https://10.1.1.10:6443
name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
    name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
user:
```

K8s Concept – Lab 1.2

YAML (YAH-ML) FILES AND PODS

- YAML is a human-readable data-serialization language used for configuration purpose
- Although `kubectl` command can do many things, in some complex scenarios using YAML file is easier for configuration management
- Commands to create a pods using `kubectl`,
 - `kubectl run nginx-kubectl --image=nginx --restart=Never`
 - `kubectl get pods -o wide` ~> to see pods information
 - `kubectl delete pods nginx-kubectl`
- Commands to create a pods using yaml,
 - `kubectl apply -f <filename>.yaml`
 - `kubectl get pods -o wide` ~> to see pods information
 - `kubectl delete pods nginx`

Example of YAML File

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
```

K8s Concept – Lab 1.3

DEPLOYMENTS

- **Deployment** defines a pod's desired behavior or characteristics.
- In layman terms, differences between Deployment and Pods are,
 - Define Replica
 - Upgrade Pods and versioning
 - Rollback and so many others
- In production, people will use Deployment rather than pods
- Commands to create a deployment,
 - **kubectl apply -f <filename>.yaml**
 - **kubectl get deployment -o wide ~>** to see deployment's information
 - **kubectl describe deployment nginx-deployment**
 - **kubectl delete deployment nginx-deployment**

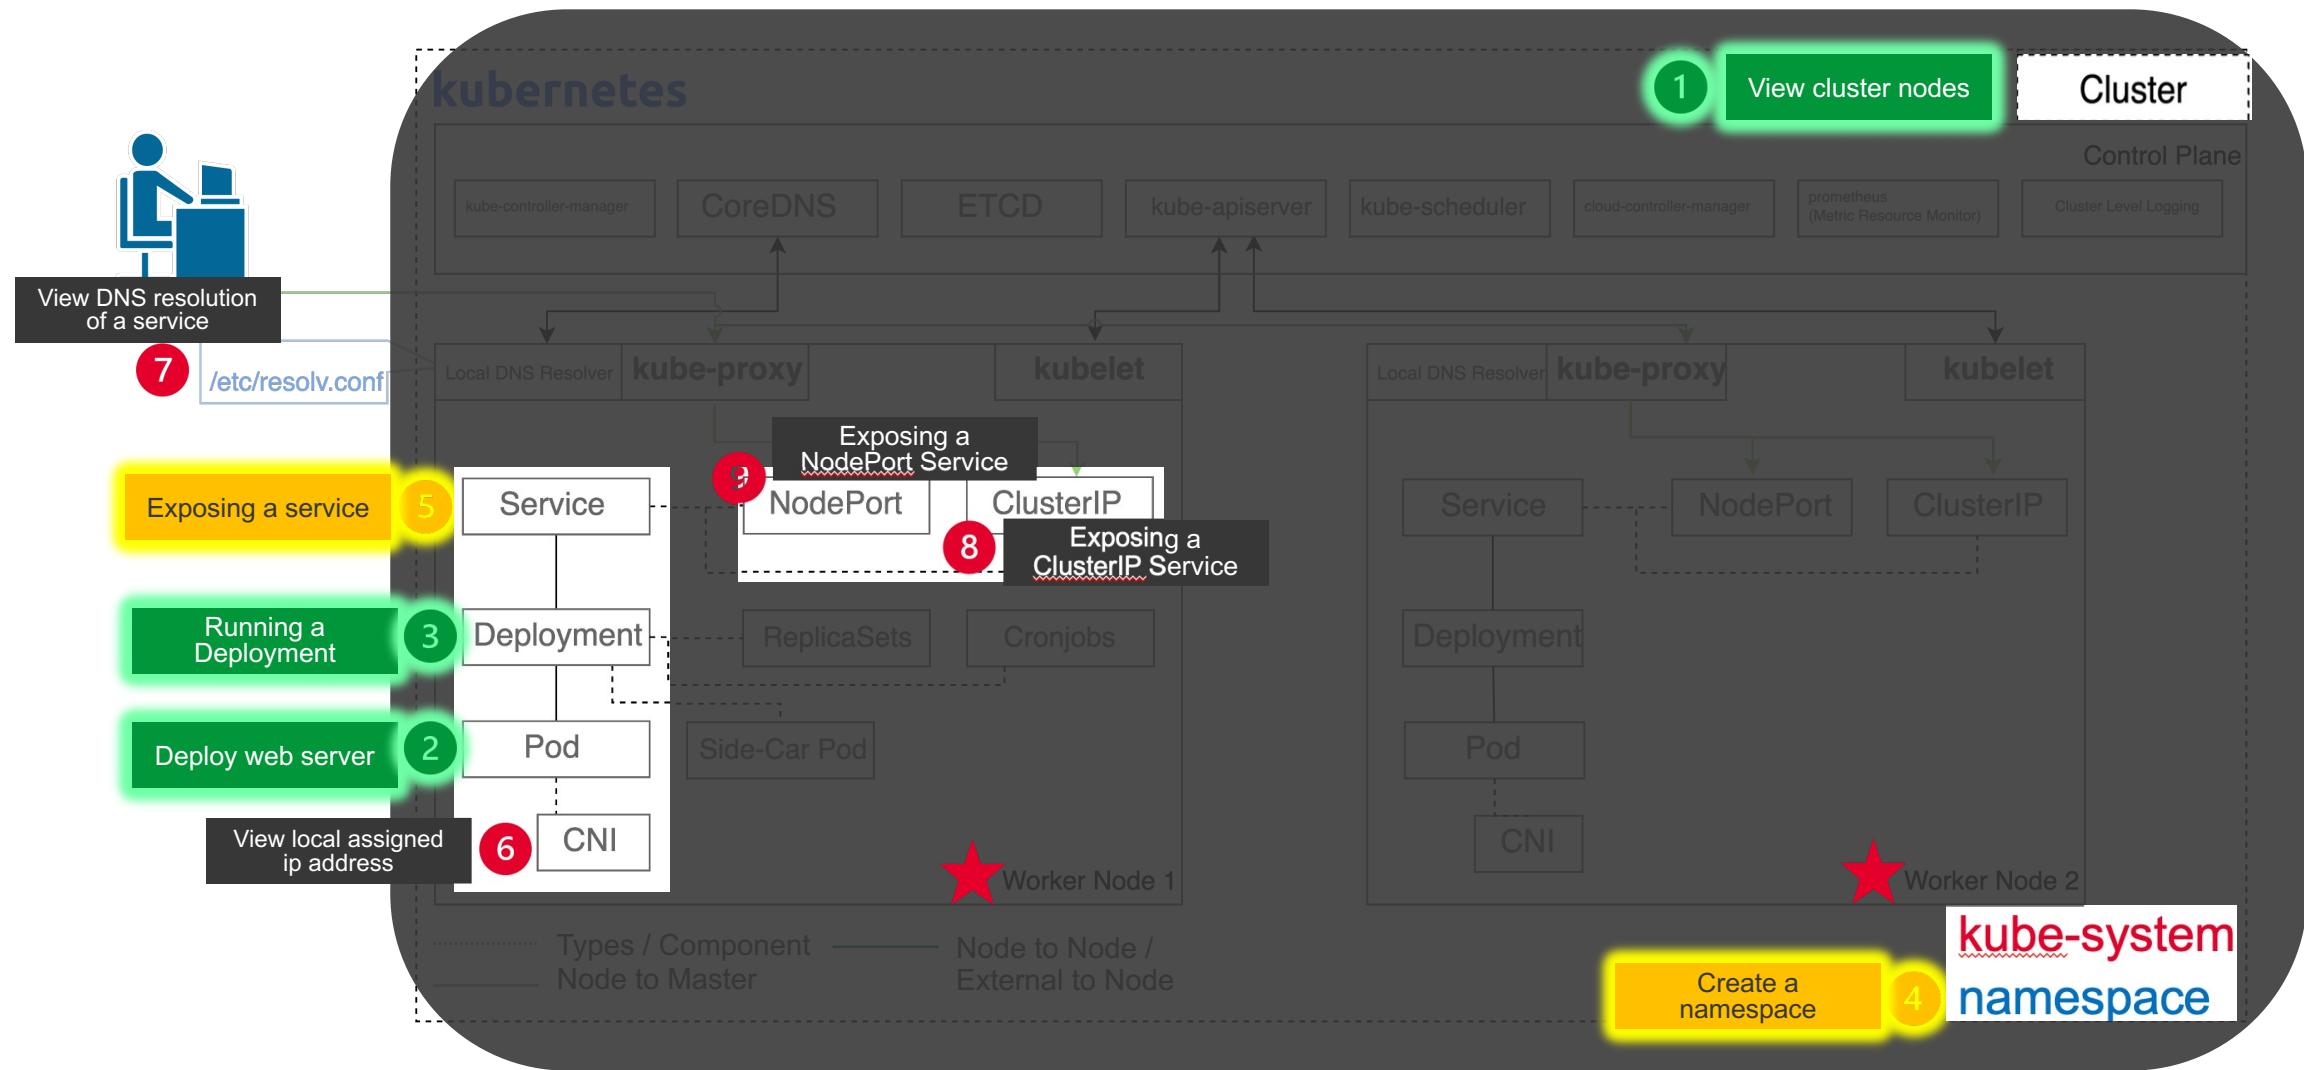
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



Lab Time (Lab 1.1 – 1.3)

20 Mins

Lab Task Overview



K8s Concept – Lab 1.4

NAMESPACES

- Namespace provides a mechanism for isolating groups of resources within a single cluster.
- There are a couple of default namespaces, <https://kubernetes.io/docs/concepts/overview/working-with-objects/namespaces/>
 - Default
 - kube-system
 - kube-public
 - kube-node-lease
- Commands to create a pod in the namespace,
 - kubectl apply -f <filename>.yaml
 - kubectl get pods
 - kubectl get pods -n production
 - kubectl describe pods mypod -n production
 - kubectl delete pods mypod -n production

```
apiVersion: v1
kind: Namespace
metadata:
  name: production
  labels:
    name: production
---
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  namespace: production
  labels:
    name: mypod
spec:
  containers:
  - name: mypod
    image: nginx
```

K8s Concept – Lab 1.5

SERVICES

- **Services** is an abstract way to expose an application running on a set of Pods as a network service
- By default, cluster IP
- Commands to create a pod and expose the service,
 - **kubectl apply -f <filename>.yaml**
 - **kubectl get pods**
 - **kubectl get service**
 - **kubectl describe pods**
 - **kubectl describe service**
 - Try to access the pods using curl from jumphost

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
---
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30080
      name: http
  selector:
    name: nginx
```



Lab Time (Lab 1.4 – 1.5)

20 Mins

Kubernetes Concepts

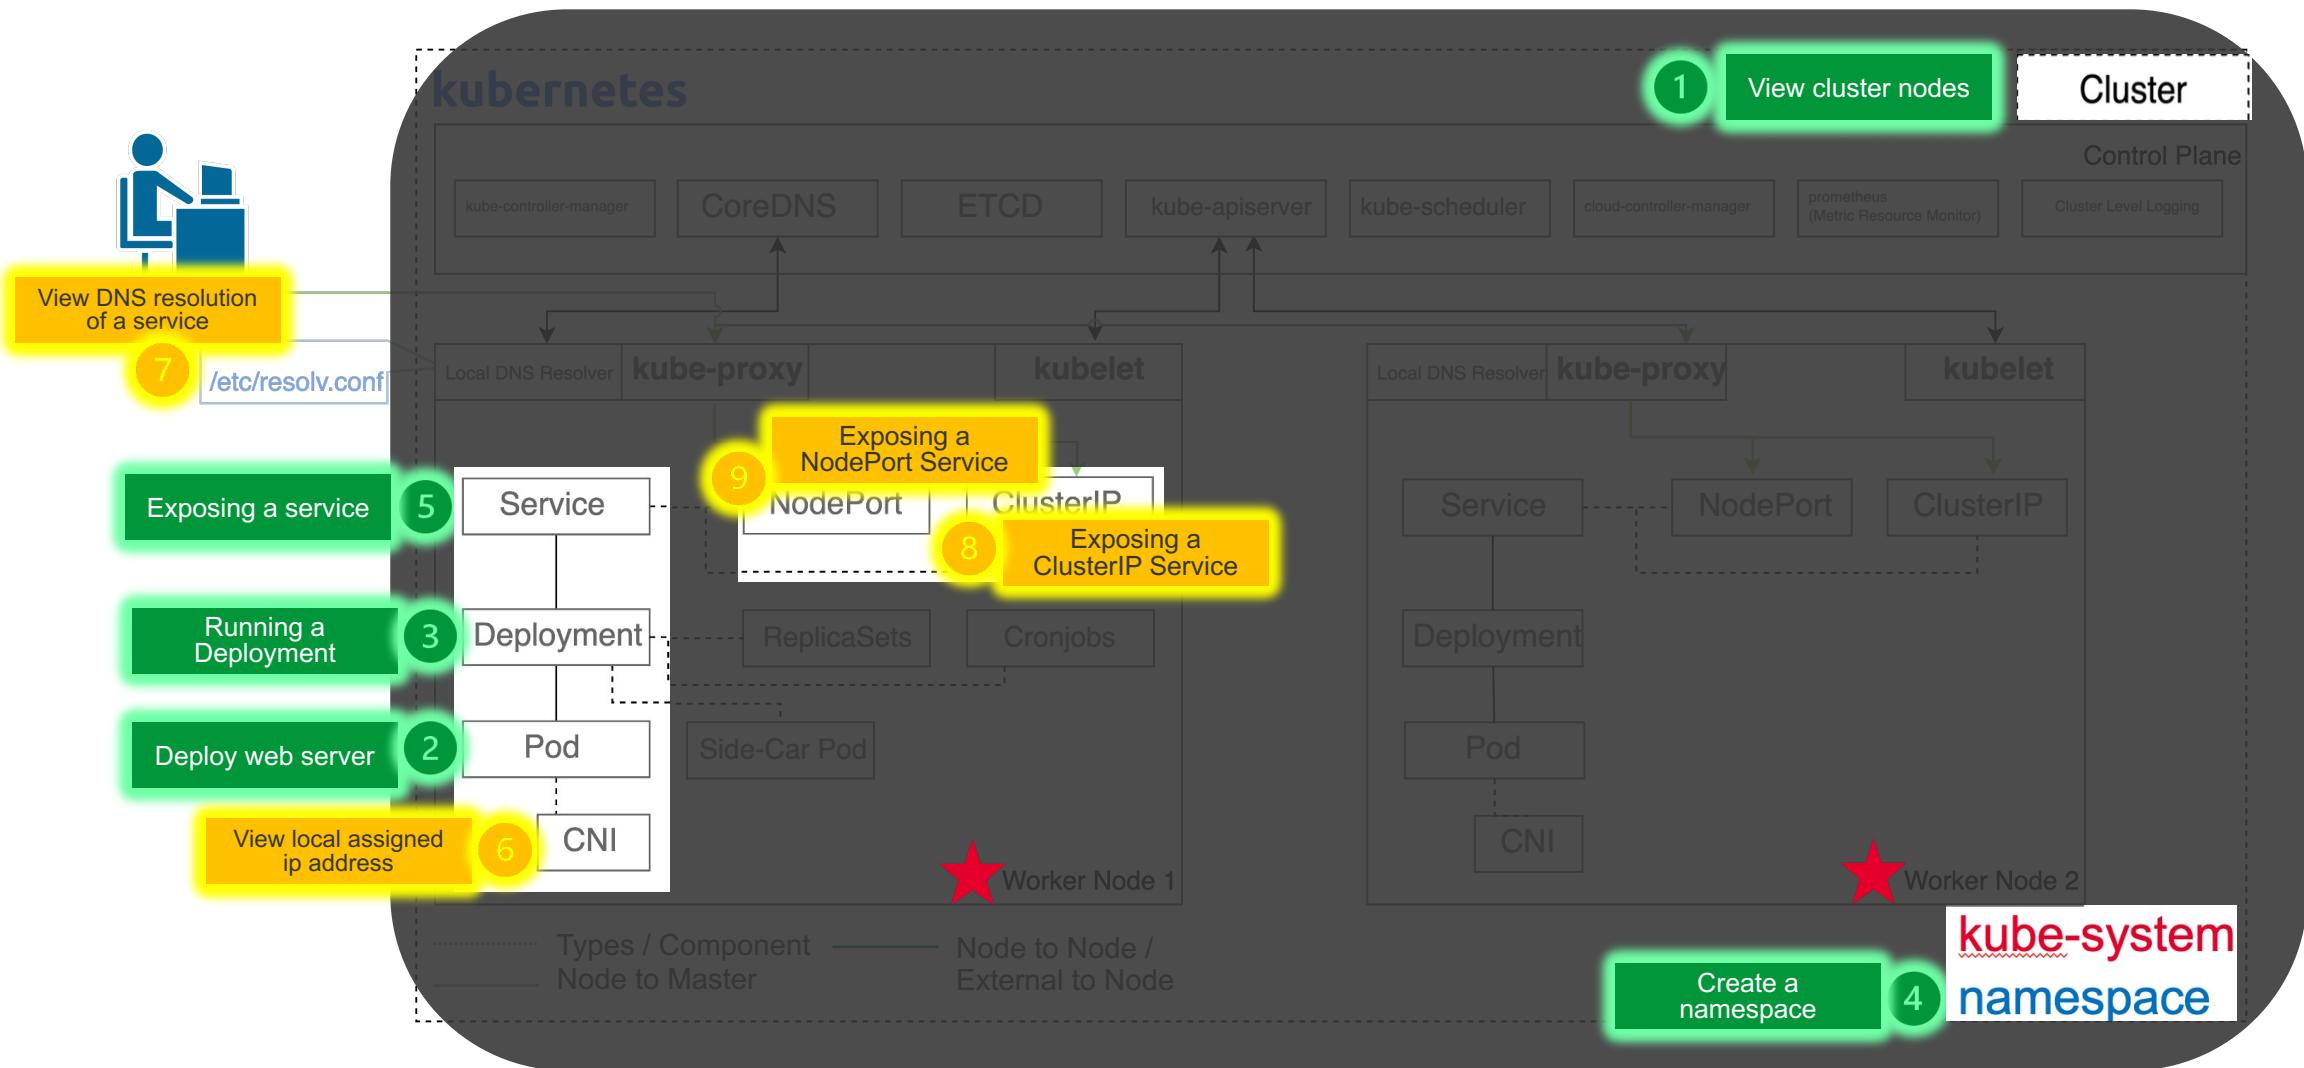
Container Network Interfaces

DNS

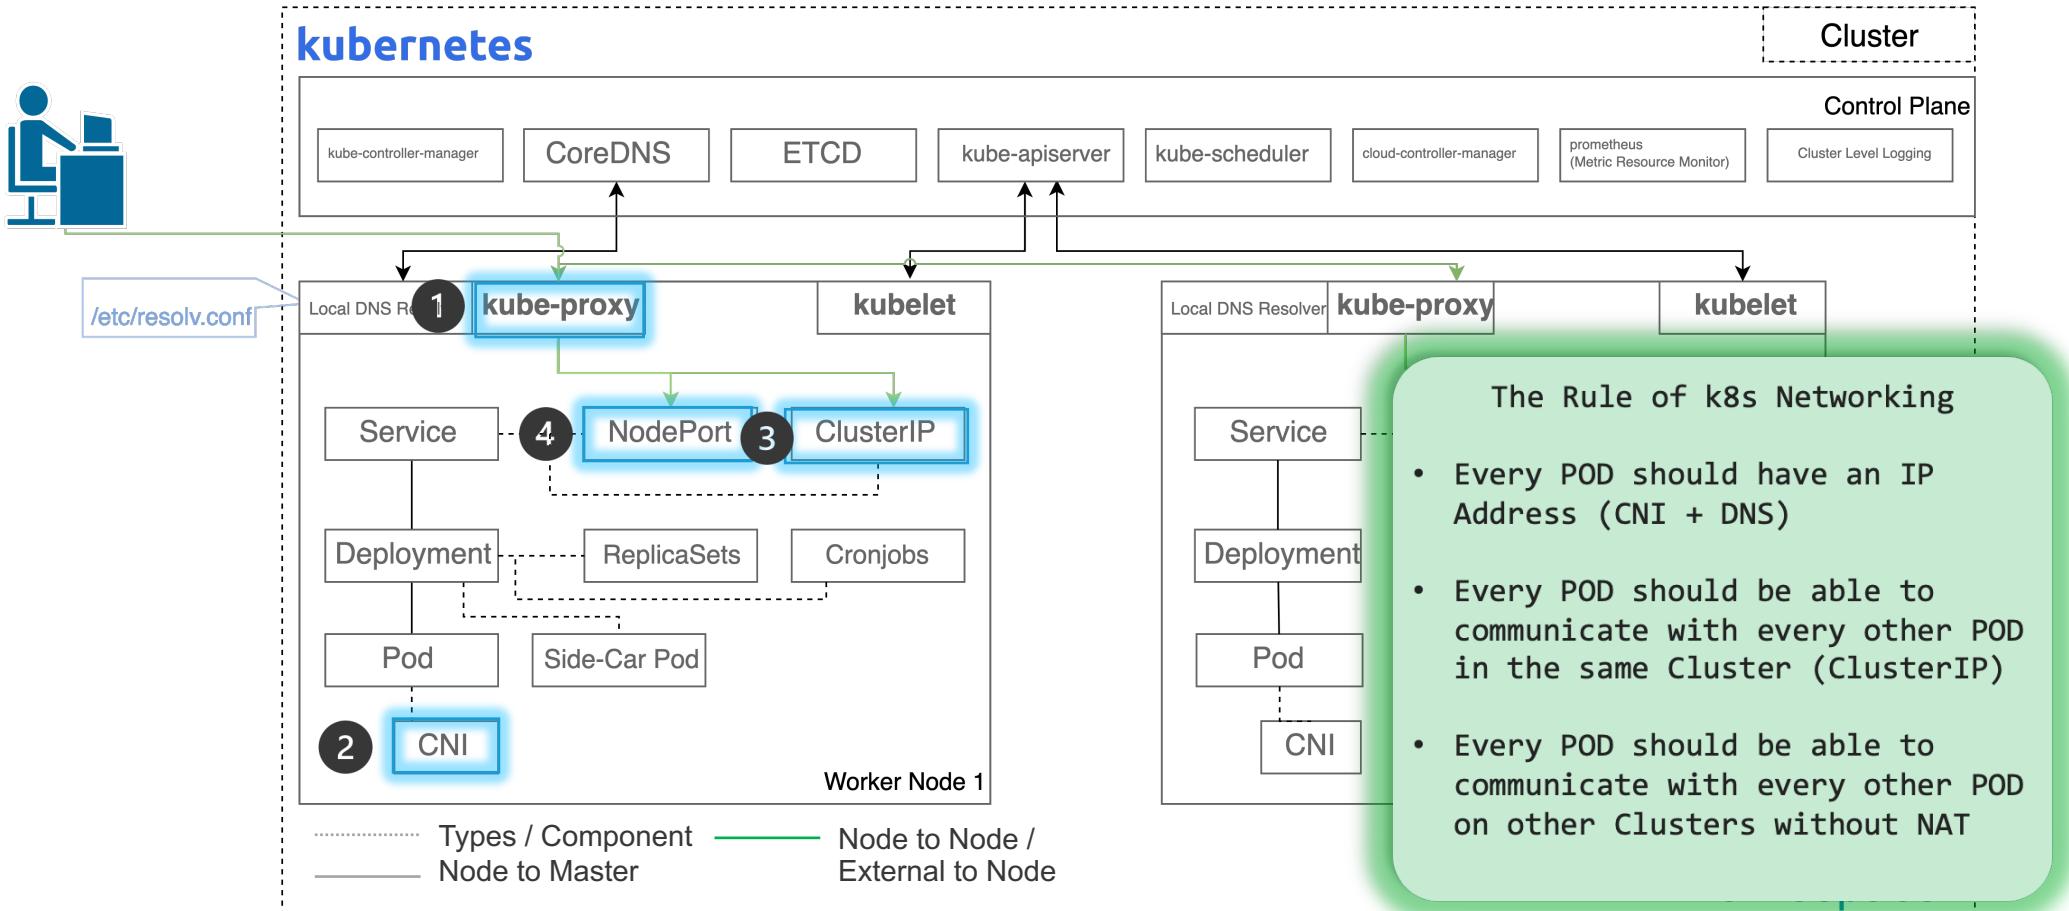
ClusterIP

NodePort

Lab Task Overview



K8s Topology: Network Connectivity



1 **kube-proxy:** a network proxy that runs on each node in your cluster. It maintains network rules on nodes.

2 **Container Network Interface:** Allows the node to be identified in the network

3 **ClusterIP:** Communication of pods in the same cluster

4 **NodePort:** Exposing Cluster to external

K8s Concept – CNI

K8 Networking

- **Container Network Interface** is a plugin that helps to configure the network, provision IP addresses, and maintain connectivity with multiple hosts
- The configuration to setup the IP subnet being used is defined in `/etc/Kubernetes/manifests/kube-controller-manager.yaml`
- Commands to be executed to know the IP address of a pods
 - **kubectl describe pods <pods' name>**
 - **kubectl get pods -o wide**

```
ubuntu@jumphost:~$ kubectl describe pods nginx-yaml
Name:           nginx-yaml
Namespace:      default
Priority:      0
Node:          worker1/10.1.1.7
Start Time:    Tue, 05 Jul 2022 02:17:24 +0000
Labels:        name=nginx-yaml
Annotations:   <none>
Status:       Running
IP:          10.244.3.65
IPs:
  IP: 10.244.3.65
```

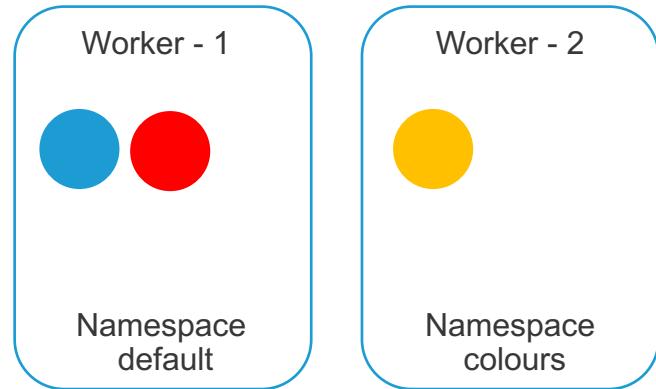
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
nginx-yaml	1/1	Running	1	7d13h	10.244.3.65	worker1
source-ip-app-86cd545f4c-q95nf	1/1	Running	2	7d23h	10.244.4.75	worker2

```
root@master1:/etc/kubernetes/manifests# cat kube-controller-manager.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    component: kube-controller-manager
    tier: control-plane
  name: kube-controller-manager
  namespace: kube-system
spec:
  containers:
  - command:
    - --experimental-cluster-signing-duration=87600h
    - kube-controller-manager
    - --allocate-node-cidrs=true
    - --authentication-kubeconfig=/etc/kubernetes/controller-manager.conf
    - --authorization-kubeconfig=/etc/kubernetes/controller-manager.conf
    - --bind-address=127.0.0.1
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --cluster-cidr=10.244.0.0/16
    - --cluster-name=kubernetes
    - --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt
    - --cluster-signing-key-file=/etc/kubernetes/pki/ca.key
    - --controllers=*,bootstrapsigner,tokencleaner
    - --kubeconfig=/etc/kubernetes/controller-manager.conf
    - --leader-elect=true
    - --port=0
    - --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt
    - --root-ca-file=/etc/kubernetes/pki/ca.crt
    - --service-account-private-key-file=/etc/kubernetes/pki/sa.key
    - --service-cluster-ip-range=10.96.0.0/12
    - --use-service-account-credentials=true
  image: k8s.gcr.io/kube-controller-manager:v1.21.1
```

K8s Concept

CORE DNS CONCEPT

- Core DNS providing a way to call the pods or services inside K8s via DNS name

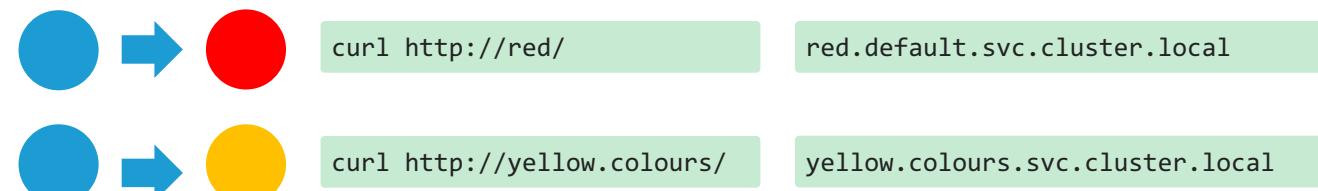


Without DNS

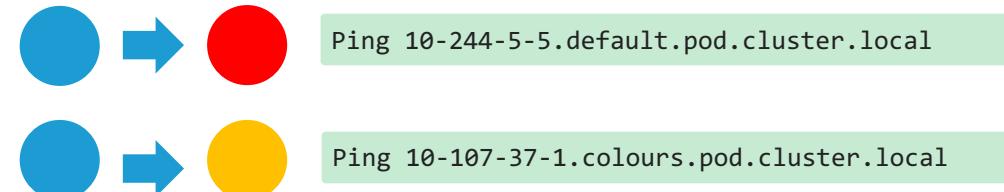


Services Name	IP address
Blue	10.244.1.5
Red	10.244.5.5
Yellow	10.107.37.1

With DNS
Calling Services



With DNS
Calling pods



The Rule of K8s Networking

- Every POD should have an IP Address
- Every POD should be able to communicate with every other POD in the same node
- Every POD should be able to communicate with every other POD on other nodes without NAT

K8s Concept

CORE DNS IS A DEFAULT MODULE (PART OF KUBE-SYSTEM)

- **Recap: Namespace**
- **Namespace** is a mechanism for isolating groups of resources within a single cluster.
- All out of the box modules will be found in **kube-system**
- Use parameter `-n kube-system`
- Use parameter `--all-namespaces`

```
kuubuntu@jumphost:~$ 
ubuntu@jumphost:~$ kubectl get namespace
NAME      STATUS   AGE
default   Active   390d
kube-node-lease Active  390d
kube-public Active   390d
kube-system Active   390d
nms       Active   77d
```

- Every POD should have an IP Address
- Every POD should be able to communicate with every other POD in the same node
- Every POD should be able to communicate with every other POD on other nodes without NAT

NAME	READY	STATUS	RESTARTS	AGE
coredns-558bd4d5db-57nt6	1/1	Running	32	390d
coredns-558bd4d5db-9h5cz	1/1	Running	36	390d
etcd-master1	1/1	Running	120	390d
etcd-master2	1/1	Running	120	390d
etcd-master3	1/1	Running	117	390d
kube-apiserver-master1	1/1	Running	159	390d
kube-apiserver-master2	1/1	Running	149	390d
kube-apiserver-master3	1/1	Running	149	390d
kube-controller-manager-master1	0/1	CrashLoopBackOff	223	16d
kube-controller-manager-master2	1/1	Running	37	390d
kube-controller-manager-master3	1/1	Running	37	390d
kube-flannel-ds-5rjn7	1/1	Running	34	390d
kube-flannel-ds-cl9pk	1/1	Running	38	390d
kube-flannel-ds-fkbv9	1/1	Running	34	390d
kube-flannel-ds-grh22	1/1	Running	33	390d
kube-flannel-ds-lrrhg	1/1	Running	38	390d
kube-proxy-5v4mg	1/1	Running	33	390d
kube-proxy-qpj5c	1/1	Running	32	390d
kube-proxy-r17fg	1/1	Running	32	390d
kube-proxy-t84x5	1/1	Running	36	390d
kube-proxy-v5pb7	1/1	Running	38	390d
kube-scheduler-master1	1/1	Running	46	390d
kube-scheduler-master2	1/1	Running	41	390d
kube-scheduler-master3	1/1	Running	36	390d
nfs-client-provisioner-7bf658486d-dk8pg	1/1	Running	30	238d

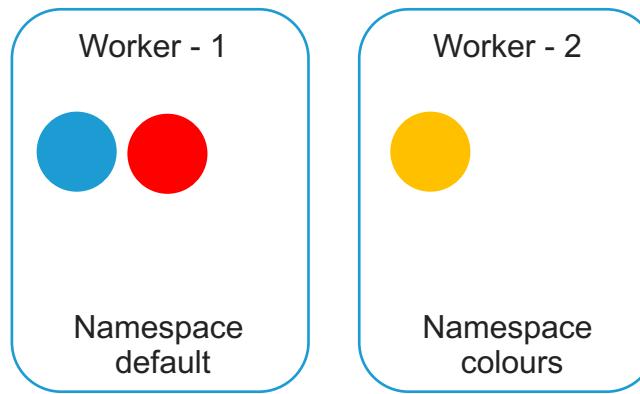
K8s Concept – DNS

CORE DNS NAMING CONVENTION

- There is a specific naming convention in K8s to call a service or pods in K8s

<Name Service/Pods>.<Namespace>.<Type(pod/svc)>.<Root Domain>

- All pods within namespace however, can call the service or pods directly without needing to call FQDN



Services Name	IP address
Blue	10.244.1.5
Red	10.244.5.5
Yellow	10.107.37.1

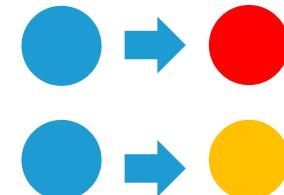
With DNS
Calling Services



curl http://red/

red.default.svc.cluster.local

With DNS
Calling pods



Ping 10-244-5-5.default.pod.cluster.local



curl http://yellow.colours/

yellow.colours.svc.cluster.local

Ping 10-107-37-1.colours.pod.cluster.local

- Commands to get into bash,
 - kubectl exec -it <pod name> -n <namespace> bash**
 - For the lab
 - Spin up pods image busybox42/alpine-pod and go into the pods
kubectl run alpine --image=busybox42/alpine-pod --restart=Never
 - Do nslookup for services that is being provisioned in previous lab
 - Bonus: Try to ping the pods that is being defined in the previous lab using DNS name

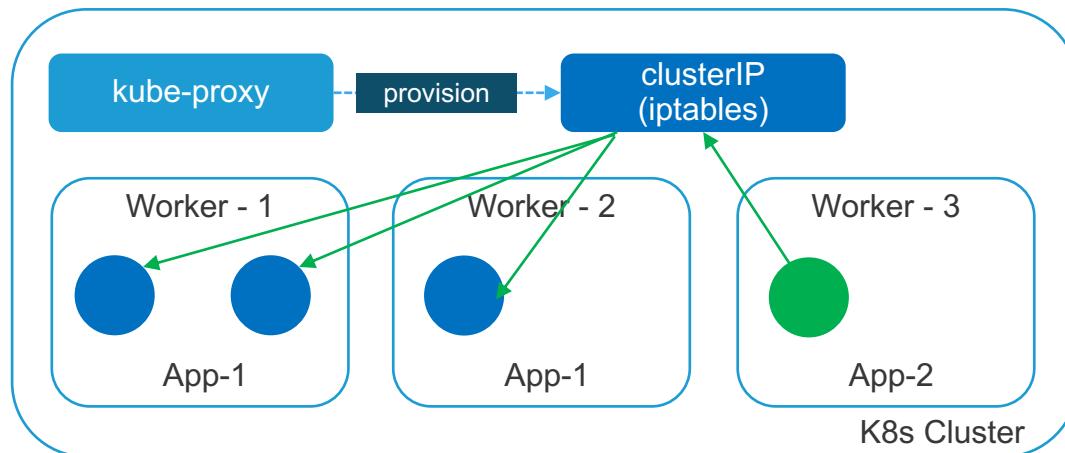
K8s Concept - ClusterIP and NodePort

EXPOSING SERVICE (CLUSTER IP AND NODEPORT) CONCEPT

Why we need service if we can go to the pods directly?

Cluster IP

Internal pods to pods communication (EW traffic)



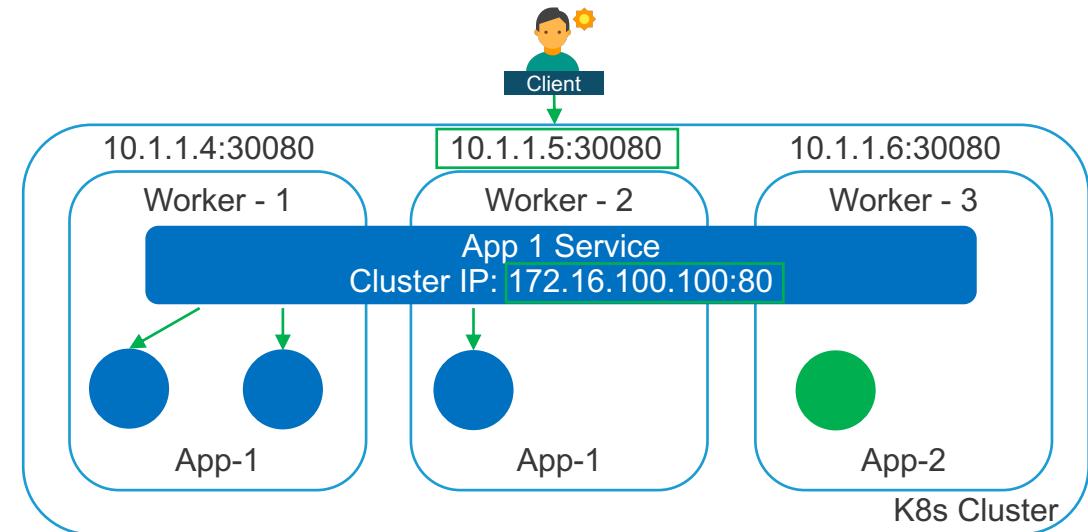
- For each Service, **kube-proxy installs iptables rules**, which capture traffic to the Service's clusterIP and port, and redirect that traffic to one of the Service's backend sets.
- By default, **kube-proxy in iptables mode chooses a backend at random**. There are different kube-proxy modes available

```
ubuntu@jumphost:~$ kubectl logs kube-proxy-5v4mg -n kube-system
I0713 08:14:12.902730 1 node.go:172] Successfully retrieved node IP: 10.1.1.8
I0713 08:14:12.903012 1 server_others.go:140] Detected node IP 10.1.1.8
W0713 08:14:12.903174 1 server_others.go:598] Unknown proxy mode "", assuming iptables proxy
I0713 08:14:12.976009 1 server_others.go:206] kube-proxy running in dual-stack mode, IPv4-primary
I0713 08:14:12.976279 1 server_others.go:212] Using iptables Proxier.
I0713 08:14:12.976329 1 server_others.go:219] creating dualStackProxier for iptables.
```

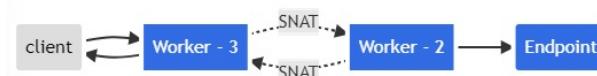
How to access my pods from outside?

NodePort

Exposing pods to outside (NS traffic)



- The **default port number range** that can be assigned in the worker node is from **30000-32767**
- The nodeport will publish the **nodeport service in all the hosts** including master nodes, regardless whether the pods is inside the host or not
- nodeport will also replace source IP Address with its own IP address

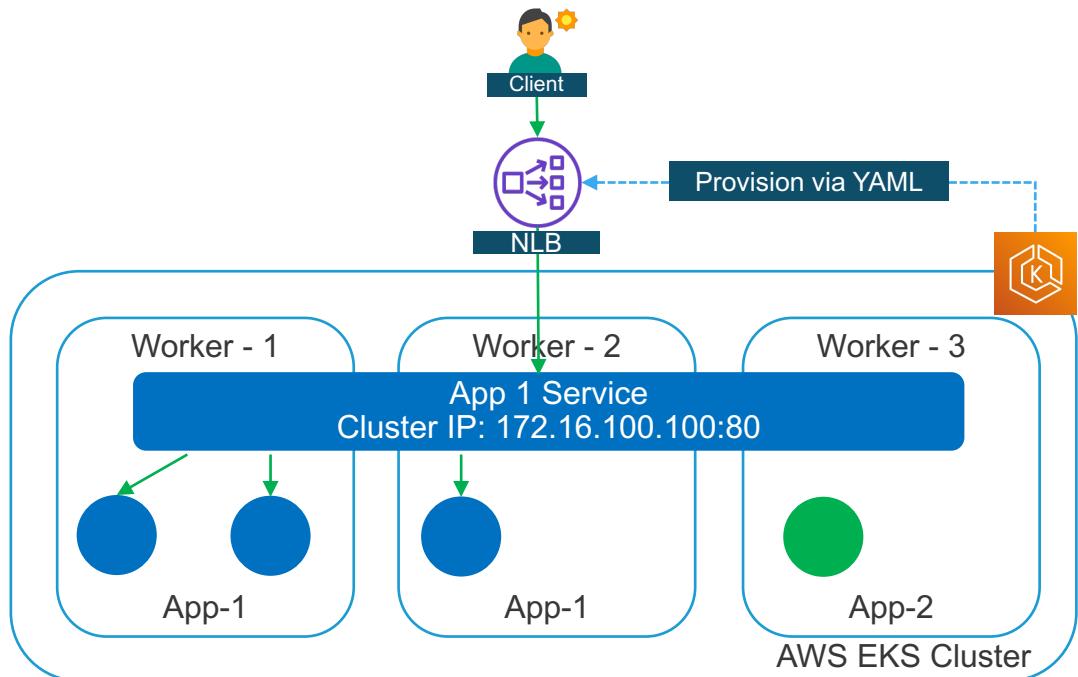


K8s Concept - LoadBalancer

LOAD-BALANCER TYPE IN PUBLIC CLOUD

How to use cloud native load balancer?

Combine both ClusterIP and NodePort



```
apiVersion: v1
kind: Service
metadata:
  name: nlb-sample-service
  namespace: nlb-sample-app
  annotations:
    service.beta.kubernetes.io/aws-load-balancer-type: external
    service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: ip
    service.beta.kubernetes.io/aws-load-balancer-scheme: internet-facing
spec:
  ports:
    - port: 80
      targetPort: 80
      protocol: TCP
  type: LoadBalancer
  selector:
    app: nginx
```



Lab Time (Lab 1.6- 1.9)

50 mins

K8s Concept

SUMMARY CNI VS KUBE-PROXY

Summary	CNI	Kube-Proxy
High Level Objectives	To create overlay networks	To intercept traffic and redirect it to correct pods
Pods provisioned on	All worker and master nodes	All worker and master nodes
Specific functions	<ul style="list-style-type: none">Configure the networkManaging IP address allocation within the clustersMaintain cluster network connectivity	<ul style="list-style-type: none">Route the traffic from pods to pods (ClusterIP)Route the traffic from outside to pods (Nodeport)Load-balancing capability for services
Issues if the function is not working	<ul style="list-style-type: none">Pods unable to get IP addressCommunication between nodes might have issues	<ul style="list-style-type: none">Pods will have issue connecting to serviceExternal will have issue reaching pods

K8s Concept

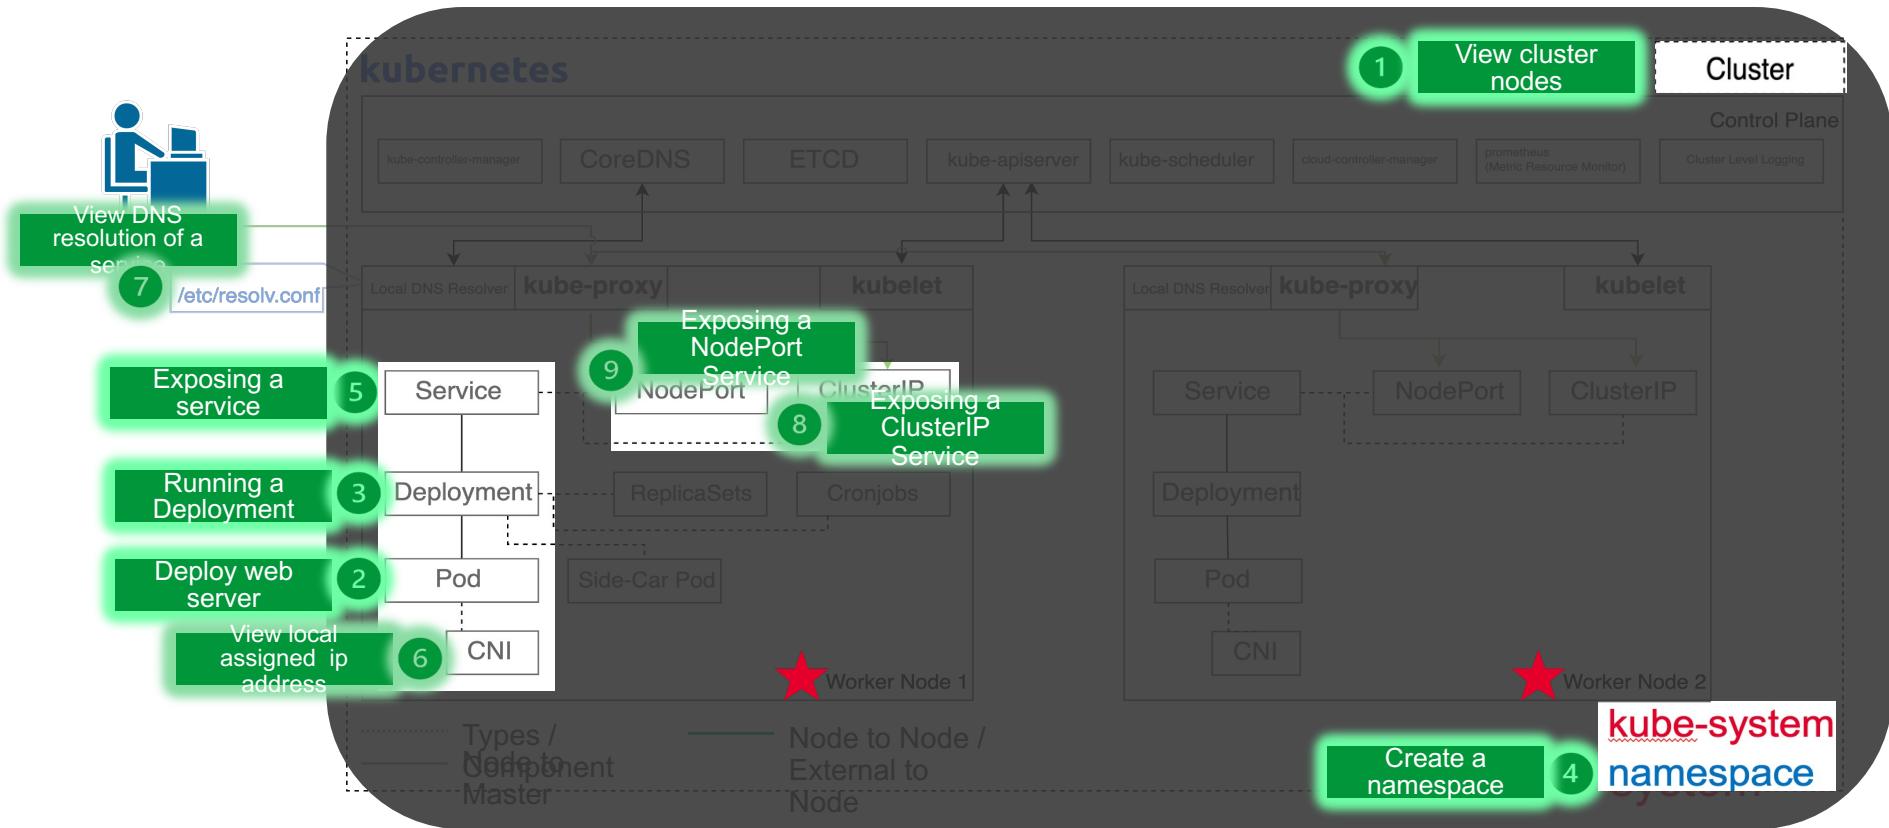
SUMMARY CLUSTER IP, NODEPORT, LOAD-BALANCER

Summary	Cluster IP	NodePort	Load-Balancer
High Level Objectives	To help connect pods to pods (internal)	To publish the port to be accessed externally	To publish the port to be accessed externally either in public VPC/VNET or private VPC/VNET
Components	Kube-proxy & CNI	Kube-proxy & CNI	Kube-proxy & CNI & Public Cloud NLB
Advantages (Plus Points)	<ul style="list-style-type: none">Simple way to publish the pods INTERNALLY within the cluster across the nodesBuilt-in Load-balancing capability to the podsCan be used for troubleshooting purpose	<ul style="list-style-type: none">Simple way to publish the pods EXTERNALLY outside of the clusterThe pods will be accessible within all nodes in the clusterBuilt-in Load-balancing capability to the pods	<ul style="list-style-type: none">Automatically provisioned NLB that will be accessible from the outsideThe NLB will load-balance the traffic directly to the pods (more efficient)The port number can be customized and not restricted to ports 30000–32767
Drawbacks (Negative Points)	<ul style="list-style-type: none">-	<ul style="list-style-type: none">No L7 routing capabilities, traffic will be passed directly to the backendRequire another Load-Balancer to LB between nodes. Otherwise, imbalance traffic might happenIt can only use ports 30000–32767If worker nodes IP address change, require changes in the surrounding environment	<ul style="list-style-type: none">No L7 routing capabilities, traffic will be passed directly to the backendAdditional price for the LB provisionedEvery service will need additional LB, it can be expensive if have many servicesOnly available on cloud environment

Quick Recap

1. K8s Concepts

- Cluster and Nodes
- Pods and Containers
- Deployments
- Service



2. K8s Networking

- Within Cluster
- Outside Cluster



Section 2 Quiz Time

10 mins



Breaktime

15 mins